

Student : Esteban Ordenes

Post Graduate Program in Data Science and Business Analytics

PGP-DSBA-UTA-Dec20-A

# Travel Package Purchase Prediction

## Context

Currently, there are 5 types of packages the company is offering - Basic, Standard, Deluxe, Super Deluxe, King. Looking at the data of the last year, we observed that 18% of the customers purchased the packages.

However, the marketing cost was quite high because customers were contacted at random without looking at the available information.

The company is now planning to launch a new product i.e. Wellness Tourism Package. Wellness Tourism is defined as Travel that allows the traveler to maintain, enhance or kick-start a healthy lifestyle, and support or increase one's sense of well-being.

However, this time company wants to harness the available data of existing and potential customers to make the marketing expenditure more efficient.

## Objective

Predict which customer is more likely to purchase the newly introduced travel package.

## Data Dictionary

- CustomerID: Unique customer ID
- ProdTaken: Whether the customer has purchased a package or not (0: No, 1: Yes)
- Age: Age of customer
- TypeofContact: How customer was contacted (Company Invited or Self Inquiry)
- CityTier: City tier depends on the development of a city, population, facilities, and living standards. The categories are ordered i.e. Tier 1 > Tier 2 > Tier 3
- Occupation: Occupation of customer
- Gender: Gender of customer
- NumberOfPersonVisiting: Total number of persons planning to take the trip with the customer
- PreferredPropertyStar: Preferred hotel property rating by customer
- MaritalStatus: Marital status of customer
- NumberOfTrips: Average number of trips in a year by customer
- Passport: The customer has a passport or not (0: No, 1: Yes)
- OwnCar: Whether the customers own a car or not (0: No, 1: Yes)

- NumberOfChildrenVisiting: Total number of children with age less than 5 planning to take the trip with the customer
- Designation: Designation of the customer in the current organization
- MonthlyIncome: Gross monthly income of the customer

## Customer interaction data:

- PitchSatisfactionScore: Sales pitch satisfaction score
- ProductPitched: Product pitched by the salesperson
- NumberOfFollowups: Total number of follow-ups has been done by the salesperson after the sales pitch
- DurationOfPitch: Duration of the pitch by a salesperson to the customer

## Loading libraries

In [163...

```
import warnings
warnings.filterwarnings("ignore")

# Libraries to help with reading and manipulating data

import pandas as pd
import numpy as np
import scipy.stats as stats

# Libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Library to split data
from sklearn.model_selection import train_test_split

# Libraries to help with model building
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree

# Removes the limit from the number of displayed columns and rows.
# This is so I can see the entire dataframe when I print it
pd.set_option("display.max_columns", None)
# pd.set_option('display.max_rows', None)
pd.set_option("display.max_rows", 200)

# To build linear model for statistical analysis and prediction
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
from statsmodels.tools.tools import add_constant

# To get different metric scores
from sklearn import metrics
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import accuracy_score, recall_score, precision_score, roc_auc_score

from sklearn.model_selection import GridSearchCV
```

```

from sklearn.model_selection import train_test_split

# For pandas profiling
#from pandas_profiling import ProfileReport

from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import GradientBoostingRegressor, AdaBoostRegressor, StackingRegressor
from sklearn.ensemble import AdaBoostClassifier, GradientBoostingClassifier
from xgboost import XGBClassifier

import warnings
warnings.filterwarnings('ignore')

```

## Read the dataset

```
In [3]: Loan = pd.read_csv("Tourism.csv")
```

```
In [4]: # copying data to another variable to avoid any changes to original data
data = Loan.copy()
```

## View the first and last 5 rows of the dataset.

```
In [5]: data.head()
```

```
Out[5]:
```

	CustomerID	ProdTaken	Age	TypeofContact	CityTier	DurationOfPitch	Occupation	Gender	Numb
0	200000	1	41.0	Self Enquiry	3	6.0	Salaried	Female	
1	200001	0	49.0	Company Invited	1	14.0	Salaried	Male	
2	200002	1	37.0	Self Enquiry	1	8.0	Free Lancer	Male	
3	200003	0	33.0	Company Invited	1	9.0	Salaried	Female	
4	200004	0	NaN	Self Enquiry	1	8.0	Small Business	Male	

```
In [6]: data.tail()
```

```
Out[6]:
```

	CustomerID	ProdTaken	Age	TypeofContact	CityTier	DurationOfPitch	Occupation	Gender	Nu
4883	204883	1	49.0	Self Enquiry	3	9.0	Small Business	Male	
4884	204884	1	28.0	Company Invited	1	31.0	Salaried	Male	
4885	204885	1	52.0	Self Enquiry	3	17.0	Salaried	Female	

	CustomerID	ProdTaken	Age	TypeofContact	CityTier	DurationOfPitch	Occupation	Gender	Nu
<b>4886</b>	204886	1	19.0	Self Enquiry	3	16.0	Small Business	Male	
<b>4887</b>	204887	1	36.0	Self Enquiry	1	14.0	Salaried	Male	

## Understand the shape of the dataset.

In [7]: `data.shape`

Out[7]: (4888, 20)

- The dataset has 4888 rows and 20 columns

## Check data types and number of non-null values for each column.

In [8]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4888 entries, 0 to 4887
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   CustomerID                           4888 non-null   int64
1   ProdTaken                            4888 non-null   int64
2   Age                                  4662 non-null   float64
3   TypeofContact                        4863 non-null   object
4   CityTier                             4888 non-null   int64
5   DurationOfPitch                      4637 non-null   float64
6   Occupation                           4888 non-null   object
7   Gender                               4888 non-null   object
8   NumberOfPersonVisiting               4888 non-null   int64
9   NumberOfFollowups                    4843 non-null   float64
10  ProductPitched                       4888 non-null   object
11  PreferredPropertyStar                 4862 non-null   float64
12  MaritalStatus                        4888 non-null   object
13  NumberOfTrips                        4748 non-null   float64
14  Passport                             4888 non-null   int64
15  PitchSatisfactionScore               4888 non-null   int64
16  OwnCar                               4888 non-null   int64
17  NumberOfChildrenVisiting             4822 non-null   float64
18  Designation                          4888 non-null   object
19  MonthlyIncome                       4655 non-null   float64
dtypes: float64(7), int64(7), object(6)
memory usage: 763.9+ KB
```

- We can see that there are total 20 columns and 4888 number of rows in the dataset.
- All columns' data type is either integer, float or object.
- There is a number of non-null values in some of the columns. We can further confirm this using `isna()` method.

In [9]: `data.isna().sum()`

Out[9]: CustomerID 0  
ProdTaken 0  
Age 226



```
0 CustomerID          4888 non-null int64
1 ProdTaken           4888 non-null int64
2 Age                 4662 non-null category
3 TypeofContact       4863 non-null category
4 CityTier            4888 non-null category
5 DurationOfPitch     4637 non-null float64
6 Occupation          4888 non-null category
7 Gender              4888 non-null category
8 NumberOfPersonVisiting 4888 non-null int64
9 NumberOfFollowups   4843 non-null float64
10 ProductPitched     4888 non-null category
11 PreferredPropertyStar 4862 non-null category
12 MaritalStatus      4888 non-null category
13 NumberOfTrips      4748 non-null float64
14 Passport           4888 non-null category
15 PitchSatisfactionScore 4888 non-null category
16 OwnCar             4888 non-null category
17 NumberOfChildrenVisiting 4822 non-null float64
18 Designation        4888 non-null category
19 MonthlyIncome      4655 non-null float64
dtypes: category(12), float64(5), int64(3)
memory usage: 366.0 KB
```

```
In [12]: # Summary of categorical columns
data[[ 'ProdTaken'
      , 'TypeofContact'
      , 'CityTier'
      , 'Occupation'
      , 'Gender'
      , 'PreferredPropertyStar'
      , 'ProductPitched'
      , 'NumberOfPersonVisiting'
      , 'NumberOfChildrenVisiting'
      , 'NumberOfFollowups'
      , 'MaritalStatus'
      , 'Designation'
      , 'Passport'
      , 'PitchSatisfactionScore'
      , 'OwnCar'
      ]].describe(include='category').T
```

Out[12]:

	count	unique	top	freq
TypeofContact	4863	2	Self Enquiry	3444
CityTier	4888	3		3190
Occupation	4888	4	Salaried	2368
Gender	4888	3	Male	2916
PreferredPropertyStar	4862	3		2993
ProductPitched	4888	5	Basic	1842
MaritalStatus	4888	4	Married	2340
Designation	4888	5	Executive	1842
Passport	4888	2		3466
PitchSatisfactionScore	4888	5		1478
OwnCar	4888	2		3032

```
In [13]: # inspect discrete columns
print('\n\nAge')
print(data.Age.value_counts())
print('\n\nTypeofContact')
print(data.TypeofContact.value_counts())
print('\n\nCityTier')
print(data.CityTier.value_counts())
print('\n\nOccupation')
print(data.Occupation.value_counts())
print('\n\nGender')
print(data.Gender.value_counts())
print('\n\nPreferredPropertyStar')
print(data.PreferredPropertyStar.value_counts())
print('\n\nProductPitched')
print(data.ProductPitched.value_counts())
print('\n\nMaritalStatus')
print(data.MaritalStatus.value_counts())
print('\n\nNumberOfPersonVisiting')
print(data.NumberOfPersonVisiting.value_counts())
print('\n\nNumberOfChildrenVisiting')
print(data.NumberOfChildrenVisiting.value_counts())
print('\n\nNumberOfFollowups')
print(data.NumberOfFollowups.value_counts())
print('\n\nDesignation')
print(data.Designation.value_counts())
print('\n\nPassport')
print(data.Passport.value_counts())
print('\n\nPitchSatisfactionScore')
print(data.PitchSatisfactionScore.value_counts())
print('\n\nOwnCar')
print(data.OwnCar.value_counts())
```

```
Age
35.0    237
36.0    231
34.0    211
31.0    203
30.0    199
32.0    197
33.0    189
37.0    185
29.0    178
38.0    176
41.0    155
39.0    150
28.0    147
40.0    146
42.0    142
27.0    138
43.0    130
46.0    121
45.0    116
26.0    106
44.0    105
51.0     90
47.0     88
50.0     86
25.0     74
52.0     68
53.0     66
48.0     65
```

49.0	65
55.0	64
54.0	61
56.0	58
24.0	56
23.0	46
22.0	46
59.0	44
21.0	41
20.0	38
19.0	32
58.0	31
57.0	29
60.0	29
18.0	14
61.0	9

Name: Age, dtype: int64

TypeofContact

Self Enquiry	3444
Company Invited	1419

Name: TypeofContact, dtype: int64

CityTier

1	3190
3	1500
2	198

Name: CityTier, dtype: int64

Occupation

Salaried	2368
Small Business	2084
Large Business	434
Free Lancer	2

Name: Occupation, dtype: int64

Gender

Male	2916
Female	1817
Fe Male	155

Name: Gender, dtype: int64

PreferredPropertyStar

3.0	2993
5.0	956
4.0	913

Name: PreferredPropertyStar, dtype: int64

ProductPitched

Basic	1842
Deluxe	1732
Standard	742
Super Deluxe	342
King	230

Name: ProductPitched, dtype: int64

MaritalStatus

Married	2340
---------	------



```
Divorced      950
Single        916
Unmarried     682
Name: MaritalStatus, dtype: int64
```

```
NumberOfPersonVisiting
3      2402
2      1418
4      1026
1         39
5         3
Name: NumberOfPersonVisiting, dtype: int64
```

```
NumberOfChildrenVisiting
1.0      2080
2.0      1335
0.0      1082
3.0       325
Name: NumberOfChildrenVisiting, dtype: int64
```

```
NumberOfFollowups
4.0      2068
3.0      1466
5.0       768
2.0       229
1.0       176
6.0       136
Name: NumberOfFollowups, dtype: int64
```

```
Designation
Executive      1842
Manager        1732
Senior Manager   742
AVP             342
VP             230
Name: Designation, dtype: int64
```

```
Passport
0      3466
1      1422
Name: Passport, dtype: int64
```

```
PitchSatisfactionScore
3      1478
5       970
1       942
4       912
2       586
Name: PitchSatisfactionScore, dtype: int64
```

```
OwnCar
1      3032
0      1856
Name: OwnCar, dtype: int64
```

- Gender seems to have rows with type Fe Male .this will be fixed

```
In [14]: def fixGenderValues(gender):  
        if gender == 'Fe Male' :  
            return 'Female'  
        else:  
            return gender
```

```
In [15]: data['Gender'] = data['Gender'].apply(fixGenderValues)
```

```
In [16]: # check for unique values for each column  
data.nunique()
```

```
Out[16]: CustomerID          4888  
ProdTaken              2  
Age                   44  
TypeofContact         2  
CityTier              3  
DurationOfPitch       34  
Occupation            4  
Gender                2  
NumberOfPersonVisiting 5  
NumberOfFollowups     6  
ProductPitched        5  
PreferredPropertyStar  3  
MaritalStatus         4  
NumberOfTrips        12  
Passport              2  
PitchSatisfactionScore 5  
OwnCar                2  
NumberOfChildrenVisiting 4  
Designation           5  
MonthlyIncome        2475  
dtype: int64
```

- We can drop 'CustomerID' column as it is an ID variable and will not add value to the model.

```
In [17]: #Dropping CustomerID columns from the dataframe  
data.drop(columns=['CustomerID'], inplace=True)
```

## Lets Evaluate the Dependant Variable - ProdTaken

```
In [18]: data['ProdTaken'].value_counts()
```

```
Out[18]: 0    3968  
        1     920  
        Name: ProdTaken, dtype: int64
```

- 18% of the customers purchased the packages

## EDA

### Univariate analysis

```
In [19]: def histogram_boxplot(feature , figsize=(15,10) , bins=None):  
        """ Histogram and Boxplot combined  
        feature: 1-d feature array  
        figsize: size of figg.default (15,10)
```

```

bins: number of bins.default None/auto
"""

mean = feature.mean()
median = feature.median()
mode = feature.mode()

f2, (ax_box2 , ax_hist2) = plt.subplots(nrows = 2, # num of rows of the subplot. gr
                                       sharex = True, # x-axis will be shared amon
                                       gridspec_kw = { "height_ratios": (.25 , .75
                                       figsize = figsize
                                       ) # create the 2 subplots

sns.boxplot(feature , ax = ax_box2 , showmeans = True , color = 'red') # boxplot wi
if bins:
    sns.distplot(feature , kde = True , ax = ax_hist2, bins = bins)
else:
    sns.distplot( feature , kde = True , ax = ax_hist2 )
ax_hist2.axvline( mean , color = 'green' , linestyle='-' , linewidth = 3 , label =
ax_hist2.axvline( median , color = 'yellow' , linestyle='-' , linewidth = 6 , label
ax_hist2.axvline( mode[0] , color = 'black' , linestyle='-' , label = 'mode' ) # ad
ax_hist2.legend()

print( 'Mean:' + str( mean ) )
print( 'Median:' + str( median ) )
print( 'Mode:' + str( mode[0] ) )

```

```

In [20]: def bar_count_pct( feature , figsize=(10,7) ):
        """
        feature : 1-d categorical feature array
        """
        mode = feature.mode()
        freq = feature.value_counts().max()

        #if isinstance(feature , int):
        #    cnt = feature.unique()
        #else:
        #    cnt = feature.unique().value_counts().sum()

        plt.figure(figsize=figsize)

        ax = sns.countplot(feature)

        total = len(feature) # Length of the column
        for p in ax.patches:
            percentage = '{:.1f}%'.format( 100 * p.get_height() / total ) # percentage of e
            x = p.get_x() + p.get_width() / 2 - 0.05 # width of the plot
            y = p.get_y() + p.get_height() # height of the plot
            ax.annotate( percentage , (x,y), size = 12) # annotate the percentage

        print( 'Top:' + str( mode[0] ) )
        print( 'Freq:' + str( freq ) )

```

## Observation on DurationOfPitch

```


In [21]: histogram_boxplot(data.DurationOfPitch)

```

```

Mean:15.490834591330602
Median:13.0
Mode:9.0

```





- DurationOfPitch feature is right-skewed.
- There are outliers to the right of the curve which may explain its skewness.

## Observation on NumberOfTrips

```
In [22]: histogram_boxplot(data.NumberOfTrips)
```

Mean:3.236520640269587  
Median:3.0  
Mode:2.0



- NumberOfTrips feature is right-skewed.
- There are outliers to the right of the curve which may explain its skewness.

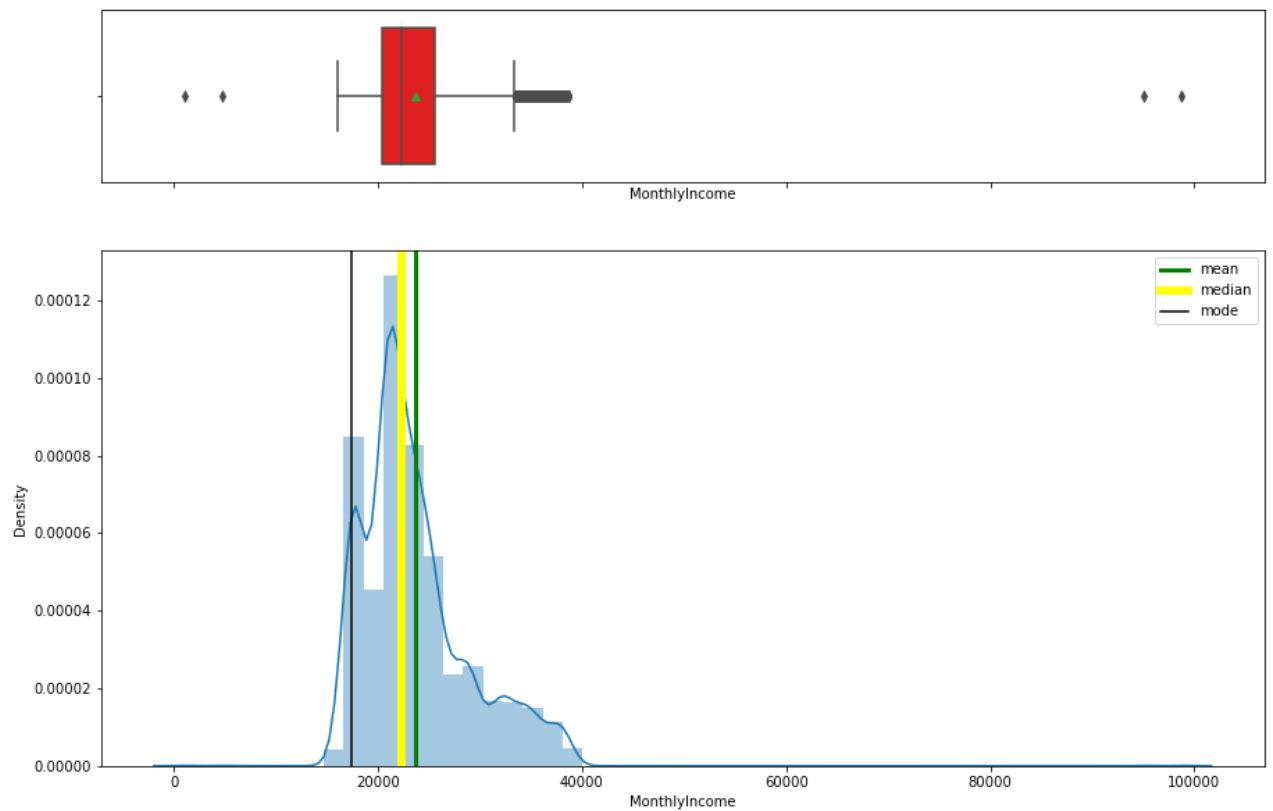
## Observation on MonthlyIncome

```
In [23]: histogram_boxplot(data.MonthlyIncome)
```

Mean:23619.85349087003

Median:22347.0

Mode:17342.0




- MonthlyIncome feature is right-skewed.
- There are outliers to the right of the curve which may explain its skewness.

## Observations on ProdTaken (Dependant Variable)

```
In [24]: print('ProdTaken\n' , data['ProdTaken'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.ProdTaken)
```

```
ProdTaken
0    0.811784
1    0.188216
Name: ProdTaken, dtype: float64
```

```
Top:0
Freq:3968
```



- Mode frequent ProdTaken is False(0) with 82%.
- Only 18% of have ProdTaken is True(1).
- There are 2 (True/False or 0/1) unique values.

## Observations on Age

```
In [25]: print('Age\n' , data['Age'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.Age)
```

```
Age
35.0    0.050837
36.0    0.049550
34.0    0.045260
31.0    0.043544
30.0    0.042686
32.0    0.042257
33.0    0.040541
37.0    0.039683
29.0    0.038181
38.0    0.037752
41.0    0.033248
39.0    0.032175
28.0    0.031532
40.0    0.031317
42.0    0.030459
27.0    0.029601
43.0    0.027885
46.0    0.025955
45.0    0.024882
26.0    0.022737
44.0    0.022523
51.0    0.019305
47.0    0.018876
```

```

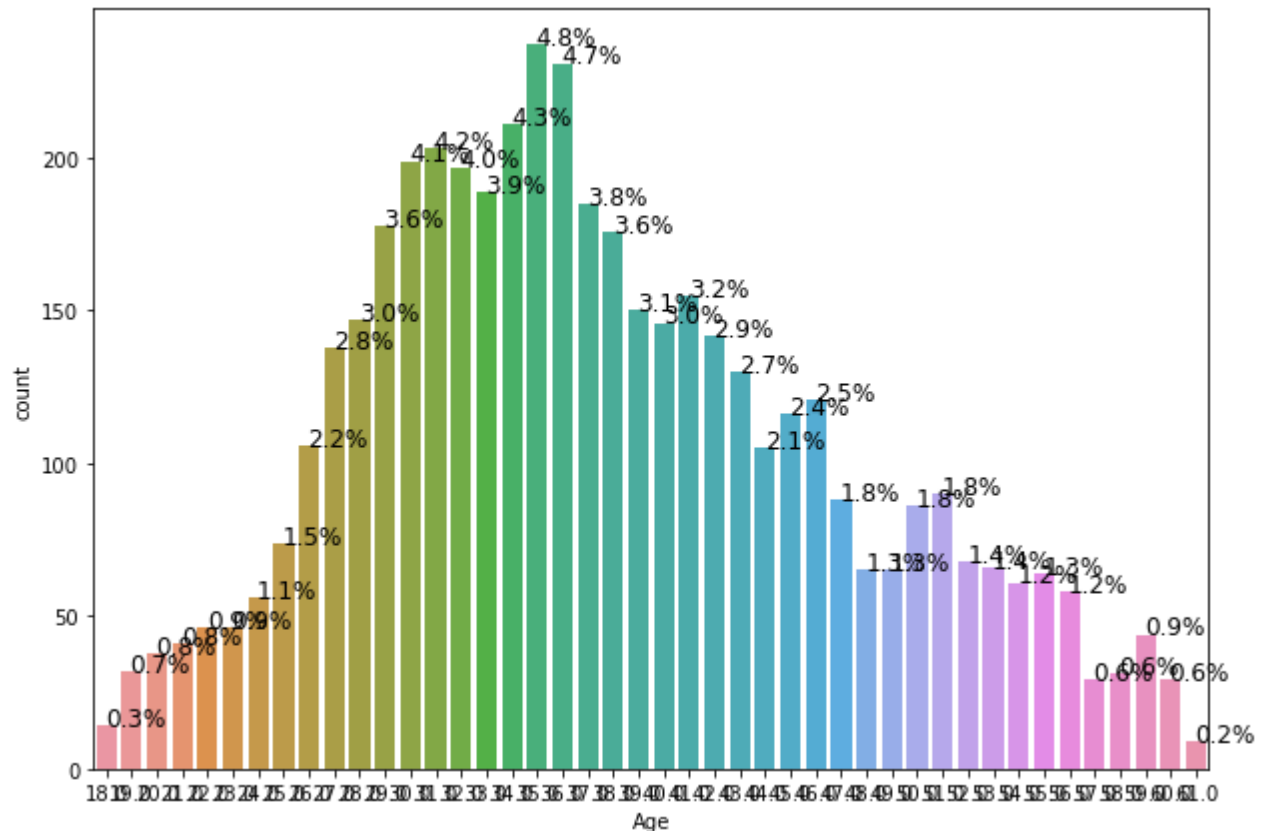
50.0    0.018447
25.0    0.015873
52.0    0.014586
53.0    0.014157
48.0    0.013943
49.0    0.013943
55.0    0.013728
54.0    0.013085
56.0    0.012441
24.0    0.012012
23.0    0.009867
22.0    0.009867
59.0    0.009438
21.0    0.008795
20.0    0.008151
19.0    0.006864
58.0    0.006650
57.0    0.006221
60.0    0.006221
18.0    0.003003
61.0    0.001931
Name: Age, dtype: float64

```

```

Top:35.0
Freq:237

```



- Most common Age is 35, it has a frequency of 237.

## Observations on TypeofContact

```

In [26]: print('TypeofContact\n' , data['TypeofContact'].value_counts(normalize=True) , '\n')
          bar_count_pct(data.TypeofContact)

```

```

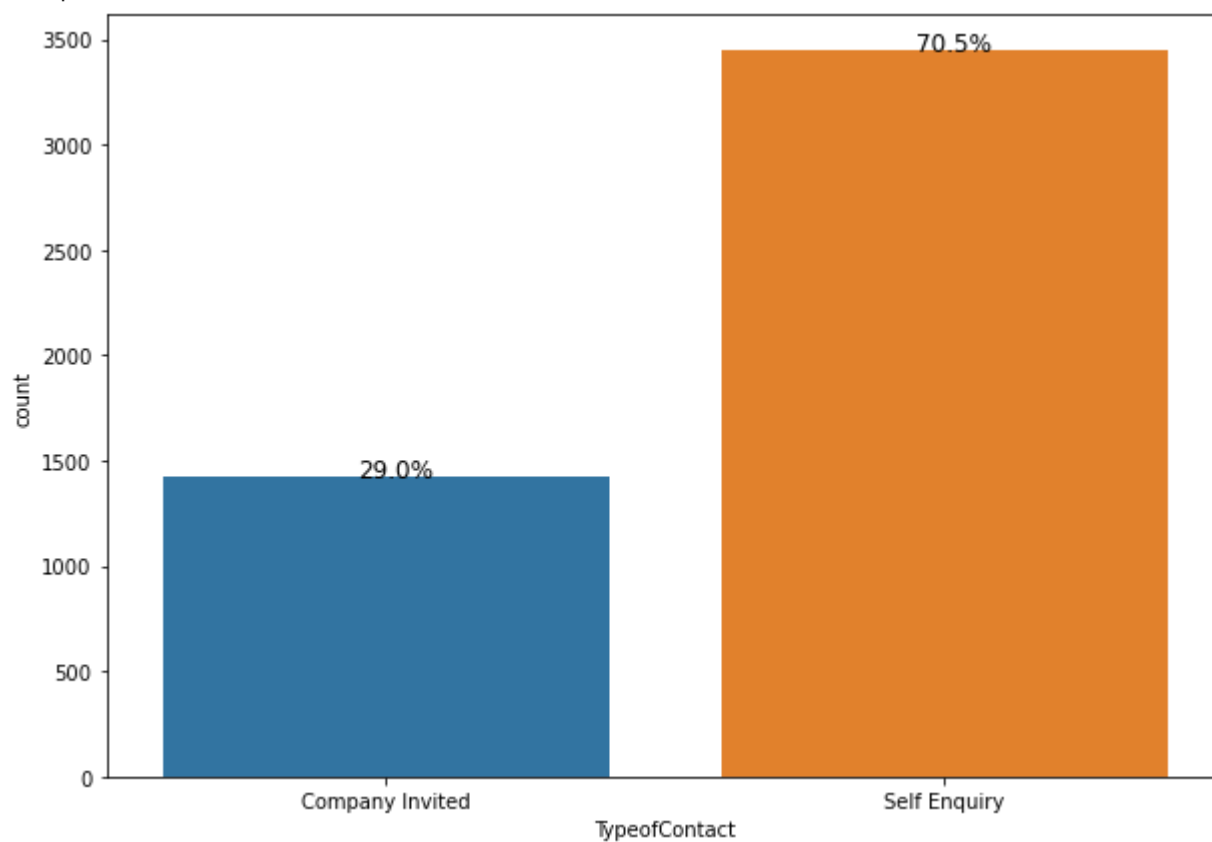
TypeofContact
Self Enquiry    0.708205

```



```
Company Invited    0.291795  
Name: TypeofContact, dtype: float64
```

```
Top:Self Enquiry  
Freq:3444
```




- Mode frequent TypeofContact is 'Self Enquiry' with 71%.
- Only 29% of have 'Company Invited'.
- There are 2 unique values.

## Observations on CityTier

```
In [27]: print('CityTier\n' , data['CityTier'].value_counts(normalize=True) , '\n')  
         bar_count_pct(data.CityTier)
```

```
CityTier  
1    0.652619  
3    0.306874  
2    0.040507  
Name: CityTier, dtype: float64
```

```
Top:1  
Freq:3190
```




- Mode frequent CityTier is Tier-1 with 65%.
- Only 4.1% of are Tier-2.
- There are 3 unique values.

## Observations on Occupation

```
In [28]: print('Occupation\n' , data['Occupation'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.Occupation)
```

```
Occupation
Salaried      0.484452
Small Business 0.426350
Large Business 0.088789
Free Lancer   0.000409
Name: Occupation, dtype: float64
```

```
Top:Salaried
Freq:2368
```




- Most frequent Occupation is Salaried with 48%.
- very few 0% are Free Lancer
- There are 4 unique values.

## Observations on Gender

```
In [29]: print('Gender\n' , data['Gender'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.Gender)
```

```
Gender
Male      0.596563
Female    0.403437
Name: Gender, dtype: float64
```

```
Top:Male
Freq:2916
```




- Most frequent Gender is Male with 60%.
- Female is 40%.
- There are 2 unique values.

## Observations on PreferredPropertyStar

```
In [30]: print('PreferredPropertyStar\n' , data['PreferredPropertyStar'].value_counts(normalize=
         bar_count_pct(data.PreferredPropertyStar))
```

```
PreferredPropertyStar
3.0    0.615590
5.0    0.196627
4.0    0.187783
Name: PreferredPropertyStar, dtype: float64
```

```
Top:3.0
Freq:2993
```




- Most frequent PreferredPropertyStar is 3 with 61%.
- 4 and 5 stars are pretty much even with 18% and 19% respectively.
- There are 3 unique values.

## Observations on ProductPitched

```
In [31]: print('ProductPitched\n' , data['ProductPitched'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.ProductPitched)
```

```
ProductPitched
Basic          0.376841
Deluxe         0.354337
Standard       0.151800
Super Deluxe   0.069967
King           0.047054
Name: ProductPitched, dtype: float64
```

```
Top:Basic
Freq:1842
```




- Most frequent ProductPitched is Basic with 38%.
- King product is the lowest with 4.7%.
- There are 5 unique values.

## Observations on NumberOfPersonVisiting

```
In [32]: print('NumberOfPersonVisiting\n' , data['NumberOfPersonVisiting'].value_counts(normaliz
          bar_count_pct(data.NumberOfPersonVisiting))
```

```
NumberOfPersonVisiting
3    0.491408
2    0.290098
4    0.209902
1    0.007979
5    0.000614
Name: NumberOfPersonVisiting, dtype: float64

Top:3
Freq:2402
```




- Most frequent NumberOfPersonVisiting is 3 with 49%.
- 5 is the lowest with 0.1%.
- There are 5 unique values.

## Observations on NumberOfChildrenVisiting

```
In [33]: print('NumberOfChildrenVisiting\n' , data['NumberOfChildrenVisiting'].value_counts(normalized=True, sort_index=True))
```

```
NumberOfChildrenVisiting
1.0    0.431356
2.0    0.276856
0.0    0.224388
3.0    0.067399
Name: NumberOfChildrenVisiting, dtype: float64
```

```
Top:1.0
Freq:2080
```



- Most frequent NumberOfChildrenVisiting is 1 with 42%.
- 3 is the lowest with 6.6%.
- There are 4 unique values.


## Observations on NumberOfFollowups

```
In [34]: print('NumberOfFollowups\n' , data['NumberOfFollowups'].value_counts(normalize=True) ,
          bar_count_pct(data.NumberOfFollowups))
```

```
NumberOfFollowups
4.0    0.427008
3.0    0.302705
5.0    0.158579
2.0    0.047285
1.0    0.036341
6.0    0.028082
Name: NumberOfFollowups, dtype: float64
```

```
Top:4.0
Freq:2068
```






- Most frequent NumberOfFollowups is 4 with 42%.
- 6 is the lowest with 2.8%.
- There are 6 unique values.

## Observations on MaritalStatus

```
In [35]: print('MaritalStatus\n' , data['MaritalStatus'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.MaritalStatus)
```

```
MaritalStatus
Married      0.478723
Divorced     0.194354
Single       0.187398
Unmarried    0.139525
Name: MaritalStatus, dtype: float64
```

```
Top:Married
Freq:2340
```



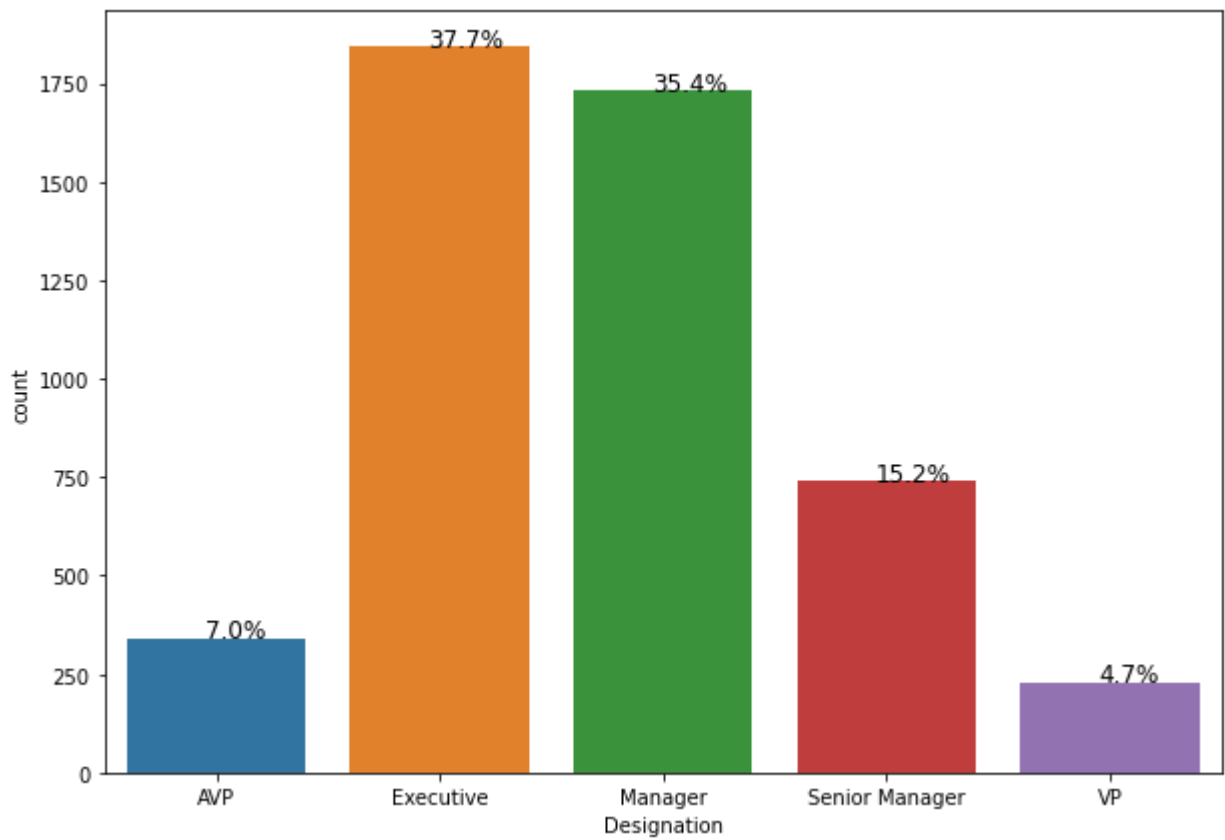
- Most frequent MaritalStatus is Married with 48%.
- Unmarried is the lowest with 14%.
- There are 4 unique values.

## Observations on Designation

```
In [36]: print('Designation\n' , data['Designation'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.Designation)
```

```
Designation
Executive      0.376841
Manager        0.354337
Senior Manager 0.151800
AVP            0.069967
VP            0.047054
Name: Designation, dtype: float64
```

```
Top:Executive
Freq:1842
```




- Most frequent Designation is Executive with 37%.
- VP is the lowest with 4.7%.
- There are 5 unique values.

## Observations on Passport

```
In [37]: print('Passport\n' , data['Passport'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.Passport)
```

```
Passport
0    0.709083
1    0.290917
Name: Passport, dtype: float64
```

```
Top:0
Freq:3466
```




- Most frequent Passport is False(0) with 71%.
- Passports True(1) is 29%.
- There are 2 unique values.

## Observations on PitchSatisfactionScore

```
In [38]: print('PitchSatisfactionScore\n' , data['PitchSatisfactionScore'].value_counts(normaliz  
bar_count_pct(data.PitchSatisfactionScore))
```

```
PitchSatisfactionScore  
3    0.302373  
5    0.198445  
1    0.192717  
4    0.186579  
2    0.119885  
Name: PitchSatisfactionScore, dtype: float64
```

```
Top:3  
Freq:1478
```




- Most frequent PitchSatisfactionScore is 3 with 30%.
- PitchSatisfactionScore 2 is the lowest with 12%.
- There are 5 unique values.

## Observations on OwnCar

```
In [39]: print('OwnCar\n' , data['OwnCar'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.OwnCar)
```

```
OwnCar
1    0.620295
0    0.379705
Name: OwnCar, dtype: float64
```

```
Top:1
Freq:3032
```



- Most frequent OwnCar is True(1) with 62%.
- OwnCar False(0) is 38%.
- There are 2 unique values.


In [ ]:

## Bivariate Analysis

In [ ]:

In [40]:

```
plt.figure(figsize=(15,5))
sns.heatmap(data.corr(),annot=True, cmap="coolwarm")
plt.show()
```



- The pairs that have a good/high correlation are:
  - NumberOfChildrenVisiting/NumberOfPersonVisiting
  - NumberOfFollowups/NumberOfPersonVisiting
  - NumberOfFollowups/NumberOfChildrenVisiting
  - NumberOfTrips/NumberOfPersonVisiting
  - NumberOfChildrenVisiting/MonthlyIncome
  - NumberOfPersonVisiting/MonthlyIncome

In [ ]:

## Bivariate analysis every possible attribute pair in relation to ProdTaken

```
In [41]: ### Function to plot distributions and Boxplots of customers
def dist_catplot(x,target='ProdTaken'):
    fig,axs = plt.subplots(2,2,figsize=(12,10))
    axs[0, 0].set_title('Distribution of a customer who take the product')
    sns.distplot(data[(data[target] == 1)][x],ax=axs[0,0],color='teal')
    axs[0, 1].set_title("Distribution of a customer who doesn't take the product")
    sns.distplot(data[(data[target] == 0)][x],ax=axs[0,1],color='orange')
    axs[1,0].set_title('Boxplot with respect to ProdTaken')
    sns.boxplot(data[target],data[x],ax=axs[1,0],palette='gist_rainbow')
    axs[1,1].set_title('Boxplot with respect to ProdTaken - Without outliers')
    sns.boxplot(data[target],data[x],ax=axs[1,1],showfliers=False,palette='gist_rainbow')
    plt.tight_layout()
    plt.show()
```

```
In [42]: ### Function to plot stacked bar charts for categorical columns
def stacked_plot(x):
    sns.set()
    ## crosstab
    tab1 = pd.crosstab(x,data['ProdTaken'],margins=True).sort_values(by=1,ascending=False)
    print(tab1)
    print('- '*120)
    ## visualising the cross tab
    tab = pd.crosstab(x,data['ProdTaken'],normalize='index').sort_values(by=1,ascending=False)
```

```

tab.plot(kind='bar',stacked=True,figsize=(17,7))
plt.legend(loc='lower left', frameon=False,)
plt.legend(loc="upper left", bbox_to_anchor=(1,1))
plt.show()


```

## Age vs ProdTaken

In [43]: `stacked_plot(data['Age'])`

ProdTaken	0	1	All
Age			
All	3786	876	4662
29.0	121	57	178
30.0	151	48	199
34.0	167	44	211
31.0	162	41	203
33.0	149	40	189
32.0	157	40	197
26.0	69	37	106
35.0	203	34	237
27.0	108	30	138
36.0	204	27	231
28.0	121	26	147
20.0	13	25	38
41.0	131	24	155
37.0	161	24	185
40.0	123	23	146
21.0	18	23	41
19.0	11	21	32
25.0	54	20	74
42.0	122	20	142
24.0	38	18	56
45.0	98	18	116
44.0	88	17	105
51.0	73	17	90
38.0	161	15	176
22.0	31	15	46
39.0	136	14	150
52.0	54	14	68
23.0	33	13	46
47.0	76	12	88
56.0	46	12	58
50.0	74	12	86
59.0	32	12	44
58.0	20	11	31
48.0	54	11	65
46.0	110	11	121
43.0	120	10	130
49.0	56	9	65
53.0	58	8	66
18.0	6	8	14
55.0	57	7	64
57.0	24	5	29
54.0	59	2	61
60.0	28	1	29
61.0	9	0	9






- Proportionally, There is significant difference in Age 18-22 for customer that take the product is a higher success rate.

## TypeofContact vs Prod Taken

In [44]: `stacked_plot(data['TypeofContact'])`

ProdTaken	0	1	All
TypeofContact			
All	3946	917	4863
Self Enquiry	2837	607	3444
Company Invited	1109	310	1419




- Proportionally, there is no significant difference in TypeofContact for customer that take the product and those who do not.

## CityTier vs Prod Taken

In [45]: `stacked_plot(data['CityTier'])`

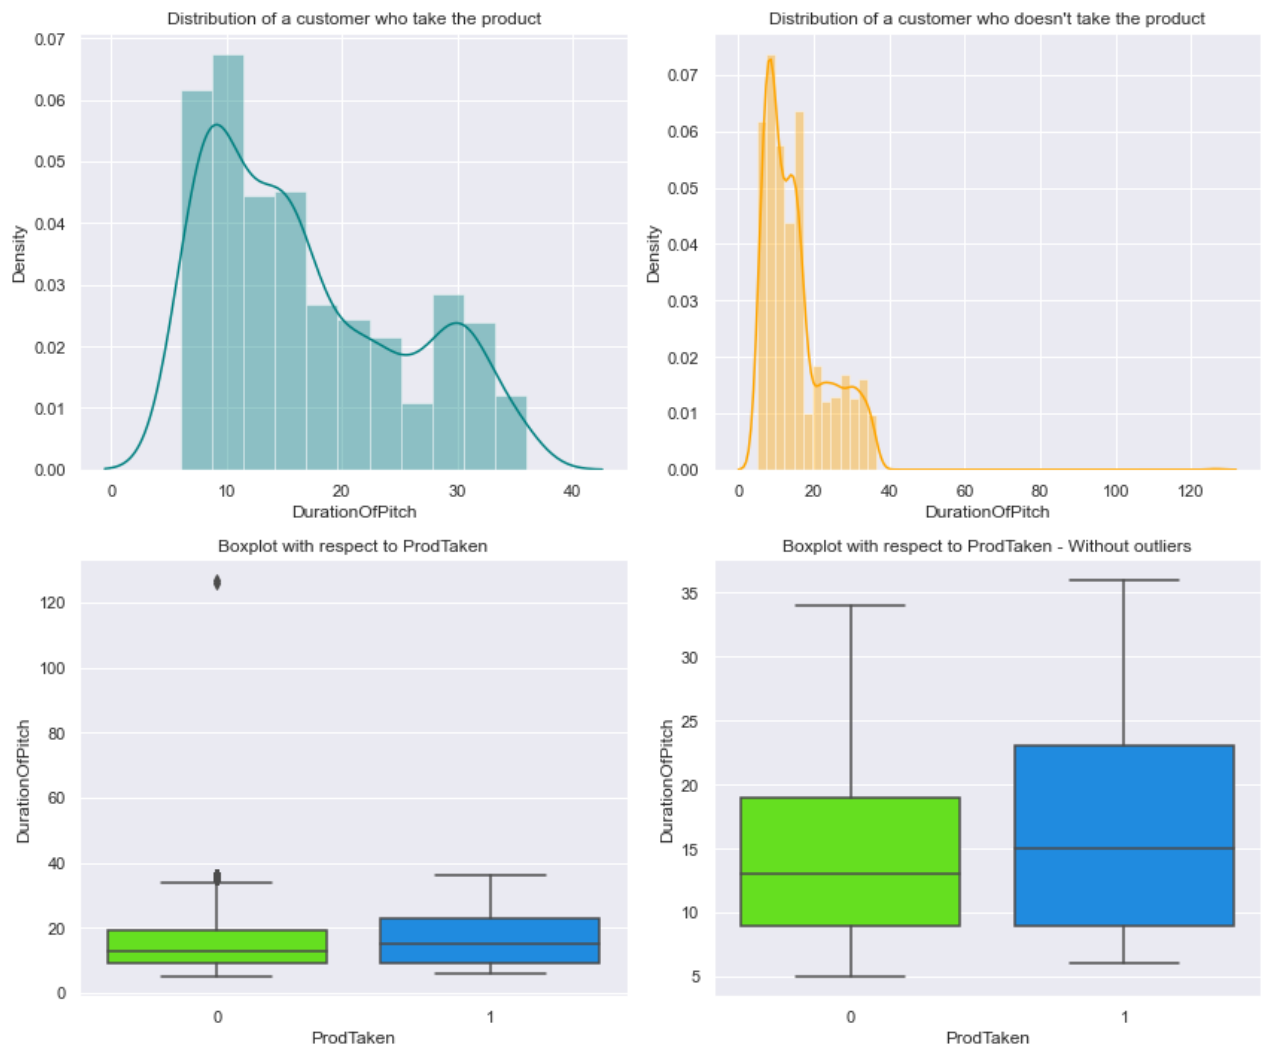
ProdTaken	0	1	All
CityTier			
All	3968	920	4888
1	2670	520	3190
3	1146	354	1500
2	152	46	198



- Proportionally, there is no significant difference in CityTier for customer that take the product and those who do not.

## DurationOfPitch vs Prod Taken

```
In [46]: dist_catplot('DurationOfPitch')
```




- When outliers are included in the comparison, there does not seem to be major difference in Duration of Pitch in regards to ProdTaken.
- When we remove outliers from the plot, we can see slight increase in ProdTaken when DurationOfPitch is higher than 20.

## Occupation vs Prod Taken

```
In [47]: stacked_plot(data['Occupation'])
```

ProdTaken	0	1	All
Occupation			
All	3968	920	4888
Salaried	1954	414	2368
Small Business	1700	384	2084
Large Business	314	120	434
Free Lancer	0	2	2




- Proportionally, there is no significant difference in Occupation between the 'Large Business', 'Small Business' and 'Salaried' categories.
- There is however a significant difference in the 'Free Lancer' category with 0%. But this category only has 2 customers which is not impactful in the overall picture.

## Gender vs Prod Taken

In [48]: `stacked_plot(data['Gender'])`

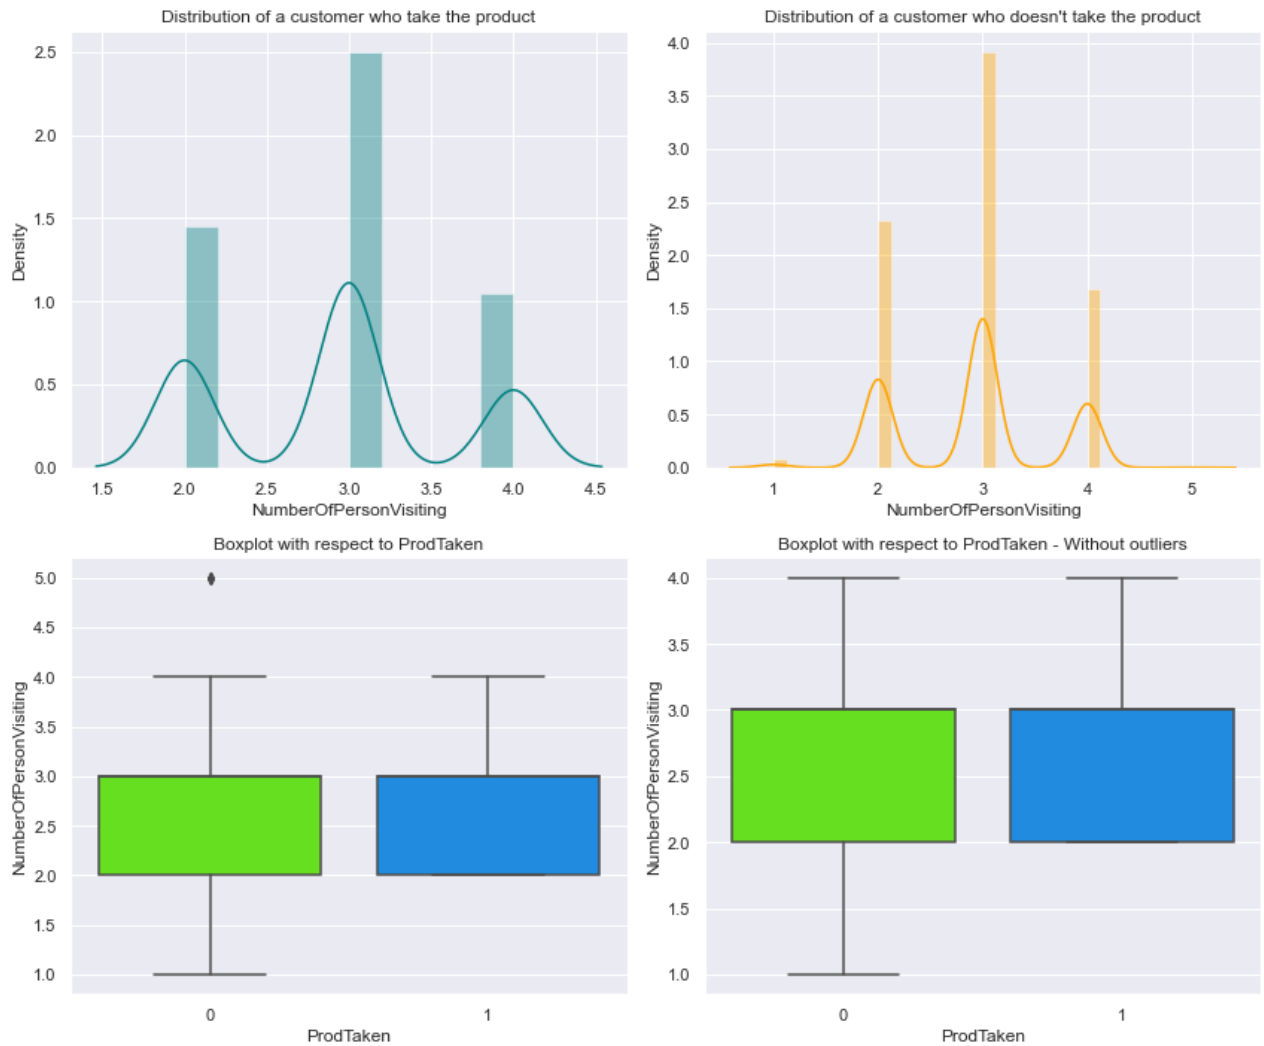
ProdTaken	0	1	All
Gender			
All	3968	920	4888
Male	2338	578	2916
Female	1630	342	1972



- Proportionally, there is no significant difference in Gender for customer that take the product and those who do not.

## NumberOfPersonVisiting vs Prod Taken

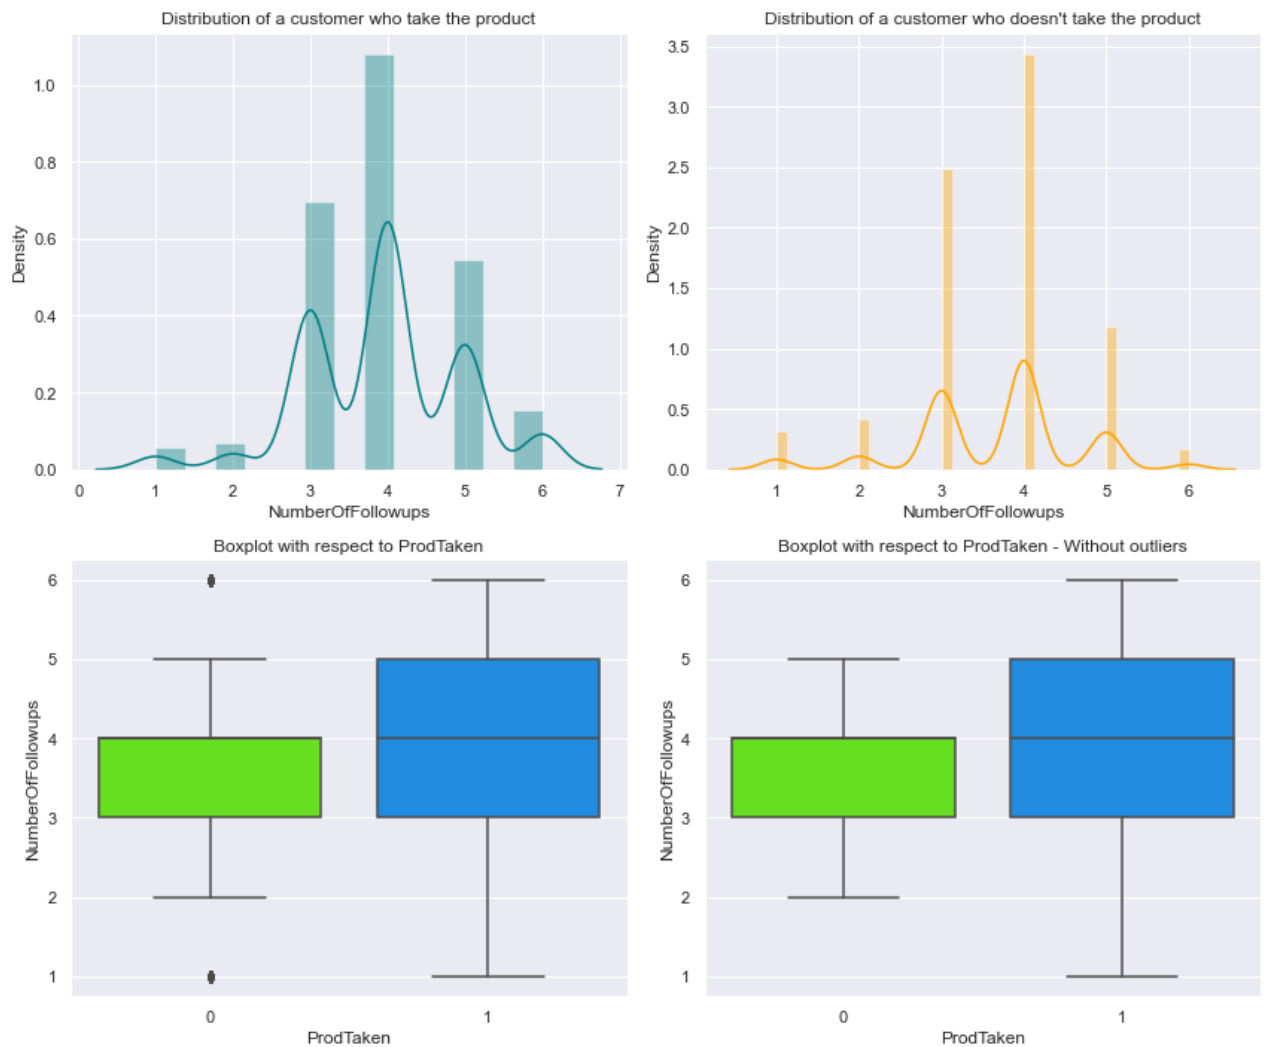
```
In [49]: dist_catplot('NumberOfPersonVisiting')
```



- There is no significant difference on `NumberOfPersonVisiting` when compared with or without outliers. In regards to `ProdTaken`.

## NumberOfFollowups vs Prod Taken

```
In [50]: dist_catplot('NumberOfFollowups')
```




- There is no significant difference on `NumberOfFollowups` when compared with or without outliers. In regards to `ProdTaken`.
- There is, however, an increase in `ProdTaken` in the ranges of 4-5 when compared to customer that did not `ProdTaken`.

## ProductPitched vs Prod Taken

```
In [51]: stacked_plot(data['ProductPitched'])
```

ProdTaken	0	1	All
ProductPitched			
All	3968	920	4888
Basic	1290	552	1842
Deluxe	1528	204	1732
Standard	618	124	742
King	210	20	230
Super Deluxe	322	20	342




- Proportionally, there is a slight increase in ProductPitched Basic. This product is also the most pitched.

## PreferredPropertyStar vs Prod Taken

In [52]: `stacked_plot(data['PreferredPropertyStar'])`

ProdTaken	0	1	All
PreferredPropertyStar			
All	3948	914	4862
3.0	2511	482	2993
5.0	706	250	956
4.0	731	182	913




- Proportionally, there is no different between all the PreferredPropertyStar categories.
- PreferredPropertyStar is the most common selected.

## MaritalStatus vs Prod Taken

In [53]: `stacked_plot(data['MaritalStatus'])`

ProdTaken	0	1	All
MaritalStatus			
All	3968	920	4888
Married	2014	326	2340
Single	612	304	916
Unmarried	516	166	682
Divorced	826	124	950






- Proportionally, Single category in MaritalStatus is the most successful.

## NumberOfTrips vs Prod Taken

```
In [54]: dist_catplot('NumberOfTrips')
```






- There is no major difference in `NumberOfTrips` in regards to `ProdTaken`.

## Passport vs Prod Taken

```
In [55]: stacked_plot(data['Passport'])
```

ProdTaken	0	1	All
Passport			
All	3968	920	4888
1	928	494	1422
0	3040	426	3466




- Customers that have `Passport` have a higher success rate on `ProdTaken`.

## PitchSatisfactionScore vs Prod Taken

In [56]: `stacked_plot(data['PitchSatisfactionScore'])`

ProdTaken	0	1	All
PitchSatisfactionScore			
All	3968	920	4888
3	1162	316	1478
5	760	210	970
4	750	162	912
1	798	144	942
2	498	88	586




- There is no major difference in `PitchSatisfactionScore` with regard to `ProdTaken`.

## OwnCar vs Prod Taken

In [57]: `stacked_plot(data['OwnCar'])`

ProdTaken	0	1	All
OwnCar			
All	3968	920	4888

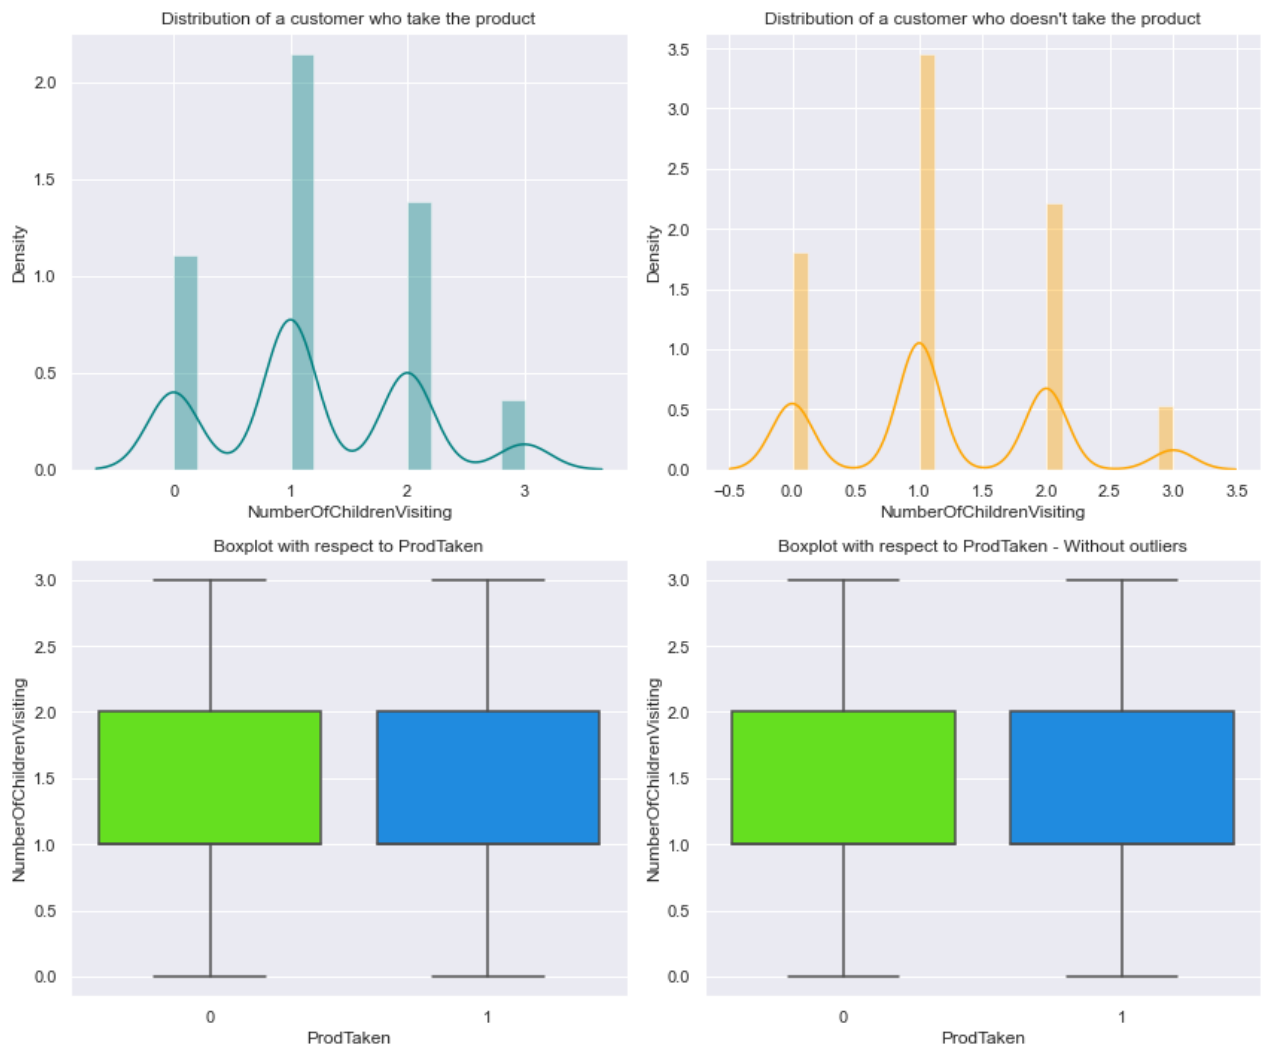
1	2472	560	3032
0	1496	360	1856



- There is no major diference in OwnCar with regard to ProdTaken.

## NumberOfChildrenVisiting vs Prod Taken

```
In [58]: dist_catplot('NumberOfChildrenVisiting')
```




- There is no major difference in `NumberOfChildrenVisiting` with regard to `ProdTaken`.

## Designation vs Prod Taken

In [59]: `stacked_plot(data['Designation'])`

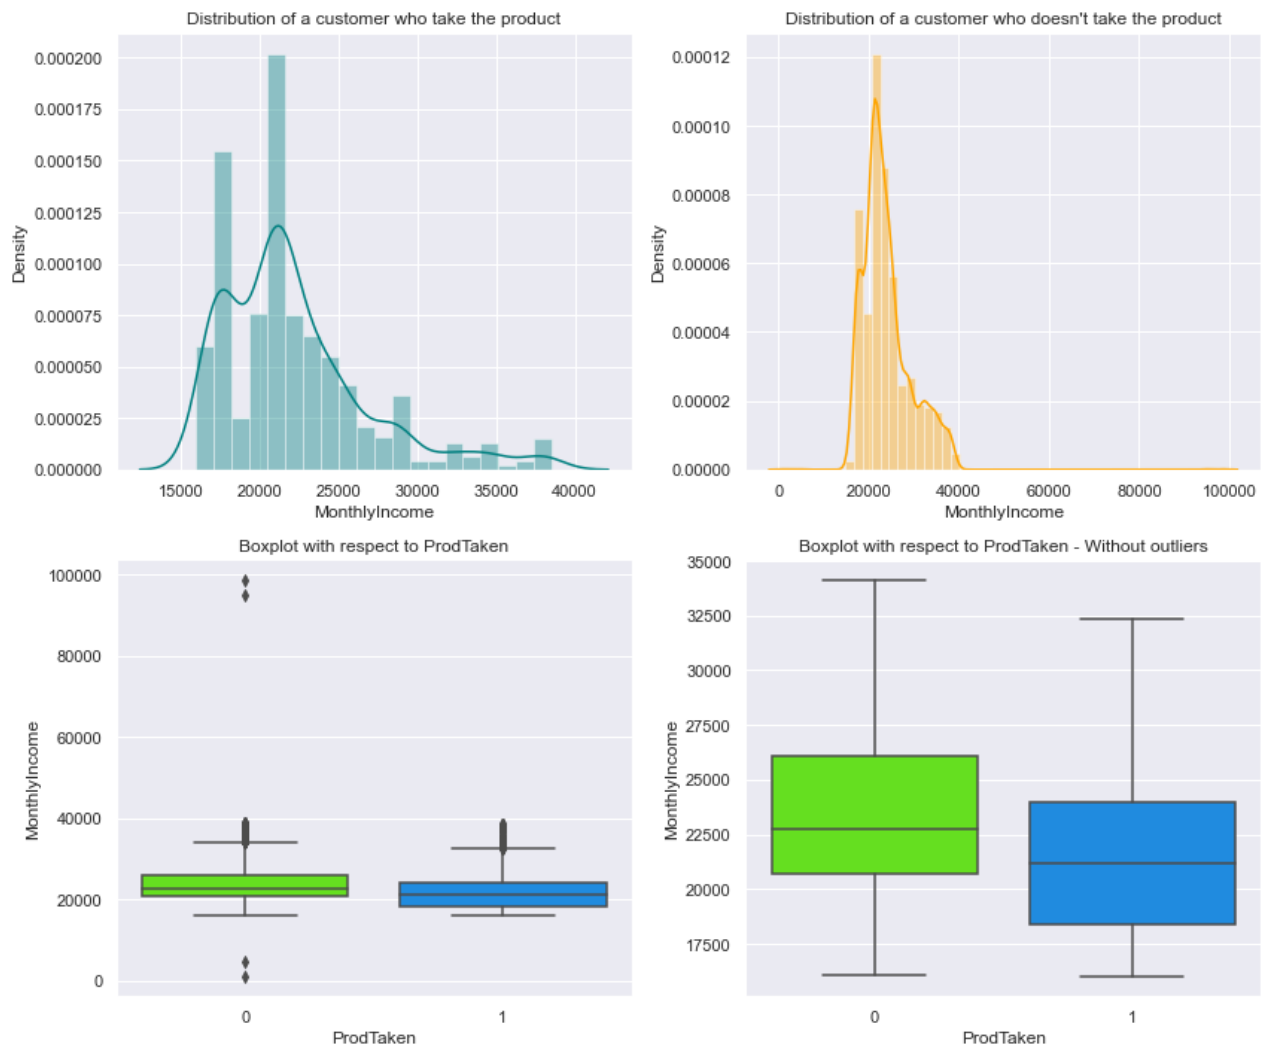
ProdTaken	0	1	All
Designation			
All	3968	920	4888
Executive	1290	552	1842
Manager	1528	204	1732
Senior Manager	618	124	742
AVP	322	20	342
VP	210	20	230



- Executive category for Designation is the most successful in regards to ProdTaken.
- All other categories seem to perform equally.

## MonthlyIncome vs Prod Taken

```
In [60]: dist_catplot('MonthlyIncome')
```



- Those customers who have an income higher than 20k-27k dollars are customers who will not take ProdTaken.
- Income seems to be a significant predictor as it provides good separation between two classes.

## Multivariate analysis

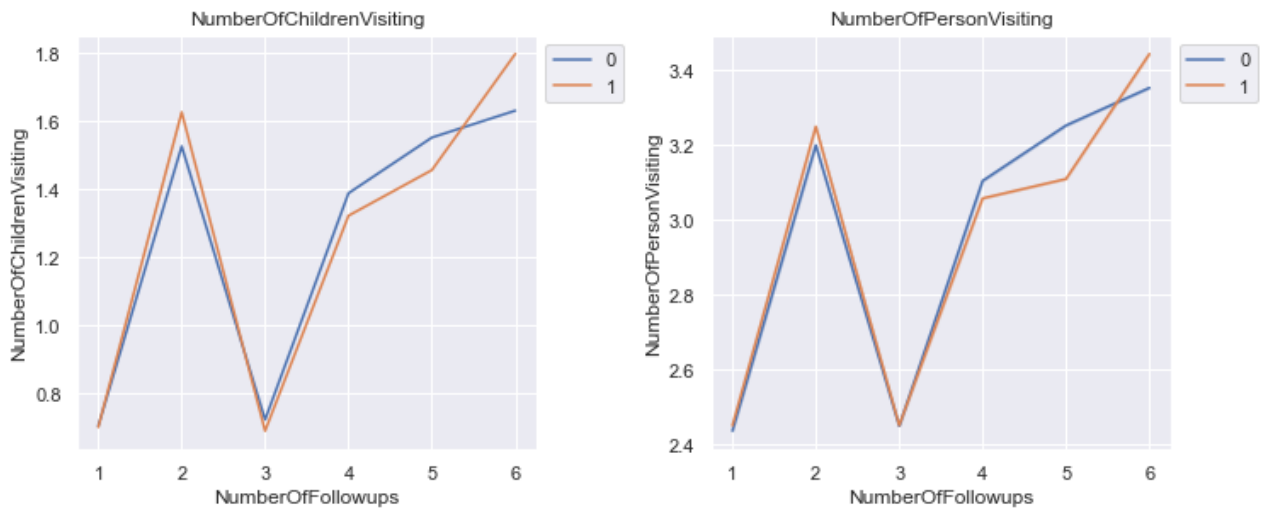
Analys variables with good and high correlation with regards to ProdTaken.

- NumberOfChildrenVisiting/NumberOfPersonVisiting
- NumberOfFollowups/NumberOfPersonVisiting
- NumberOfFollowups/NumberOfChildrenVisiting
- NumberOfTrips/NumberOfPersonVisiting
- NumberOfChildrenVisiting/MonthlyIncome
- NumberOfPersonVisiting/MonthlyIncome

```
In [61]: cols = data[['NumberOfChildrenVisiting', 'NumberOfPersonVisiting']].columns.tolist()
plt.figure(figsize=(15,12))
for i, variable in enumerate(cols):
    plt.subplot(3,3,i+1)
    sns.lineplot(data['NumberOfFollowups'], data[variable], hue=data['Pr
```

```
plt.tight_layout()
plt.title(variable)
plt.legend(bbox_to_anchor=(1, 1))


plt.show()
```



- When visualizing difference between NumberOfChildrenVisiting and NumberOfPersonVisiting in regards to NumberOfFollowups and Prod Taken, there really is no significant difference. This could be a sign of Multicollinearity.

```
In [62]: cols = data[['NumberOfChildrenVisiting', 'MonthlyIncome']].columns.tolist()
plt.figure(figsize=(15,12))
for i, variable in enumerate(cols):
    plt.subplot(3,3,i+1)
    sns.lineplot(data['NumberOfTrips'], data[variable], hue=data['ProdTa
    plt.tight_layout()
    plt.title(variable)
    plt.legend(bbox_to_anchor=(1, 1))

plt.show()
```




When visualizing difference between NumberOfChildrenVisiting and NumberOfPersonVisiting in regards to NumberOfTrips and Prod Taken, there seems to be a major difference on the extreme(outlier) cases. This may not be influential, otherwise this may be a case of multicollinearity.

```
In [63]: cols = data[['NumberOfPersonVisiting', 'MonthlyIncome']].columns.tolist()
```

```
plt.figure(figsize=(15,12))
for i, variable in enumerate(cols):
    plt.subplot(3,3,i+1)
    sns.lineplot(data['NumberOfTrips'],data[variable],hue=data['ProdTa
    plt.tight_layout()
    plt.title(variable)
    plt.legend(bbox_to_anchor=(1, 1))

plt.show()
```



- There is a difference between NumberOfPersonsVisiting and MonthlyIncome, in regards to NumberOfTrips and ProdTaken.

In [ ]:

## Data Pre-processing

### Null treatment

- Out of 4888 rows, the following columns have null values

In [64]: `data.isnull().sum()`

```
Out[64]: ProdTaken          0
Age          226
TypeofContact  25
CityTier      0
DurationOfPitch 251
Occupation     0
Gender         0
NumberOfPersonVisiting  0
NumberOfFollowups  45
ProductPitched  0
PreferredPropertyStar  26
MaritalStatus   0
NumberOfTrips   140
Passport        0
PitchSatisfactionScore  0
OwnCar          0
NumberOfChildrenVisiting 66
Designation     0
```



```
MonthlyIncome      233  
dtype: int64
```


## Treat Age nulls

```
In [65]: data.Age.unique()
```

```
Out[65]: [41.0, 49.0, 37.0, 33.0, NaN, ..., 52.0, 47.0, 18.0, 60.0, 61.0]  
Length: 45  
Categories (44, float64): [41.0, 49.0, 37.0, 33.0, ..., 47.0, 18.0, 60.0, 61.0]
```

```
In [66]: plt.figure(figsize=(15, 7))  
sns.countplot(y="Age", data=data, order=data["Age"].value_counts().index[1:30])
```

```
Out[66]: <AxesSubplot:xlabel='count', ylabel='Age'>
```



- There are 45 difference ages, we will try to reduce this and group by AgeRanges.
- First, we must treat null values and outliers
- We will treat null values by replacing these with the Mode.

```
In [67]: data.Age.mode()
```

```
Out[67]: 0    35.0  
Name: Age, dtype: category  
Categories (44, float64): [18.0, 19.0, 20.0, 21.0, ..., 58.0, 59.0, 60.0, 61.0]
```


```
In [68]: data['Age'] = data['Age'].fillna(data['Age'].mode()[0])
```

## Treat TypeofContact nulls

```
In [69]: data['TypeofContact'] = data['TypeofContact'].cat.add_categories('Unknown')  
data['TypeofContact'].fillna('Unknown', inplace=True)
```

```
In [70]: bar_count_pct(data.TypeofContact)
```

```
Top:Self Enquiry  
Freq:3444
```



### Treat DurationOfPitch nulls

```
In [71]: data['DurationOfPitch'] = data['DurationOfPitch'].fillna(data['DurationOfPitch'].mean())
```

### Treat NumberOfFollowups nulls

```
In [72]: data['NumberOfFollowups'] = data['NumberOfFollowups'].fillna(data['NumberOfFollowups'].mean())
```

### Treat PreferredPropertyStar nulls

```
In [73]: data['PreferredPropertyStar'] = data['PreferredPropertyStar'].fillna(data['PreferredPropertyStar'].mean())
```

### Treat NumberOfTrips nulls

```
In [74]: data['NumberOfTrips'] = data['NumberOfTrips'].fillna(data['NumberOfTrips'].mode()[0])
```

### Treat NumberOfChildrenVisiting nulls

```
In [75]: data['NumberOfChildrenVisiting'] = data['NumberOfChildrenVisiting'].fillna(data['NumberOfChildrenVisiting'].mean())
```

### Treat MonthlyIncome nulls

```
In [76]: data['MonthlyIncome'] = data['MonthlyIncome'].fillna(data['MonthlyIncome'].mean())
```

```
In [ ]:
```

```
In [77]: data.isnull().sum()
```

```
Out[77]: ProdTaken      0
```

```

Age                                0
TypeofContact                      0
CityTier                           0
DurationOfPitch                     0
Occupation                         0
Gender                             0
NumberOfPersonVisiting              0
NumberOfFollowups                   0
ProductPitched                     0
PreferredPropertyStar               0
MaritalStatus                      0
NumberOfTrips                       0
Passport                           0
PitchSatisfactionScore              0
OwnCar                             0
NumberOfChildrenVisiting            0
Designation                        0
MonthlyIncome                      0
dtype: int64

```

- There are no null values now

## Outlier treatment

```

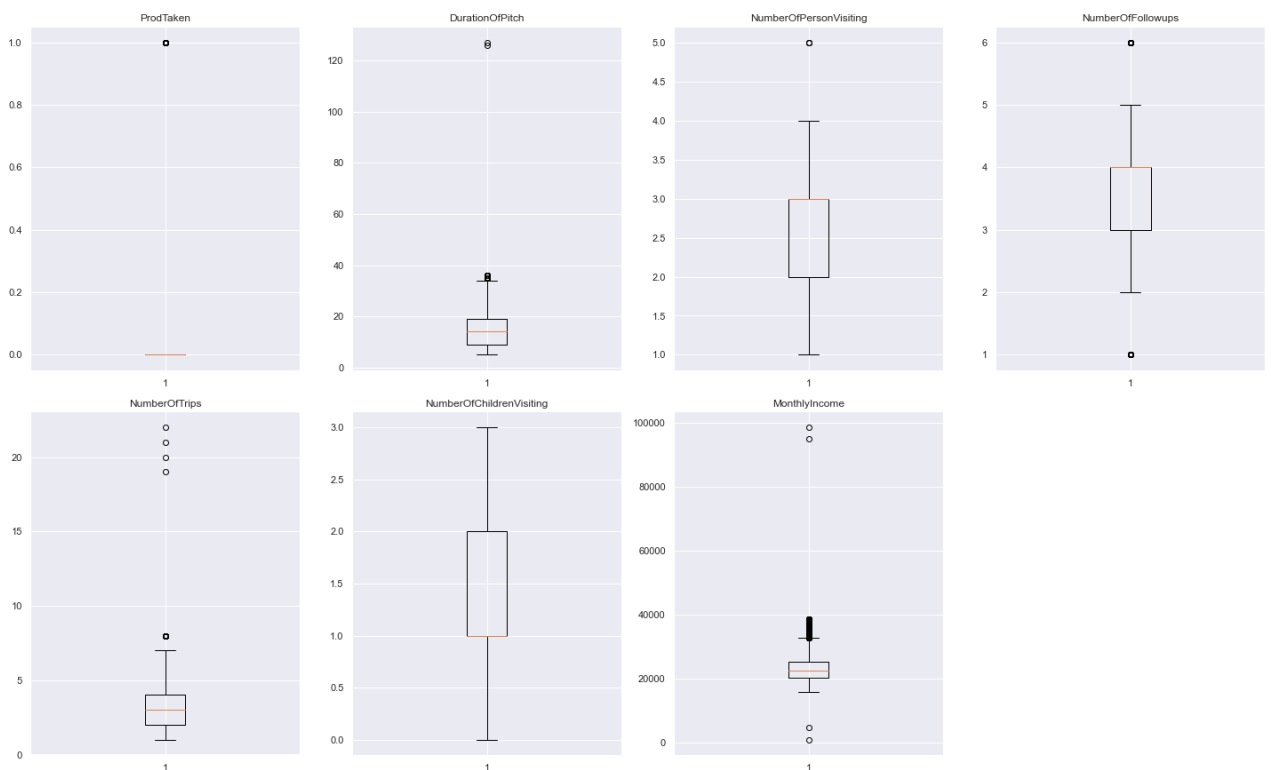
In [78]: numerical_col = data.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(20,30))

for i, variable in enumerate(numerical_col):
    plt.subplot(5,4,i+1)
    plt.boxplot(data[variable],whis=1.5)
    plt.tight_layout()
    plt.title(variable)

plt.show()

```



- We will remove `ProdTaken` from outlier treatment since this is a boolean, and also is the dependant variable.

```
In [79]: def treat_outliers(df,col):
    """
    treats outliers in a variable
    col: str, name of the numerical variable
    df: data frame
    col: name of the column
    """
    Q1=df[col].quantile(0.25) # 25th quantile
    Q3=df[col].quantile(0.75) # 75th quantile
    IQR=Q3-Q1
    Lower_Whisker = Q1 - 1.5*IQR
    Upper_Whisker = Q3 + 1.5*IQR
    df[col] = np.clip(df[col], Lower_Whisker, Upper_Whisker) # all the values smaller t
                                                            # and all the values above

    return df

def treat_outliers_all(df, col_list):
    """
    treat outlier in all numerical variables
    col_list: list of numerical variables
    df: data frame
    """
    for c in col_list:
        df = treat_outliers(df,c)

    return df
```

```
In [80]: numerical_col = data.select_dtypes(include=np.number).columns.tolist()# getting list of

# items to be removed
unwanted= {'ProdTaken'} # these column have very few non zero observation , doing outli

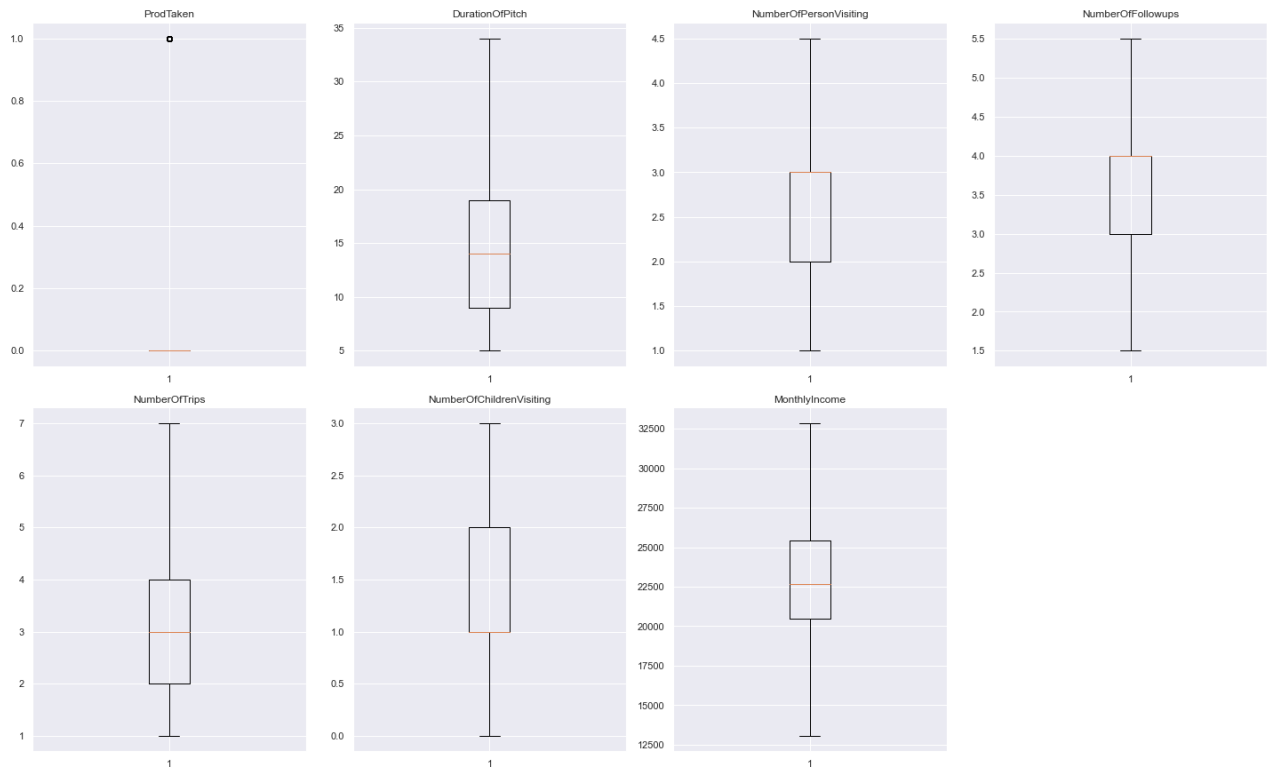
numerical_col = [ele for ele in numerical_col if ele not in unwanted]
data = treat_outliers_all(data,numerical_col)
```

```
In [81]: numerical_col = data.select_dtypes(include=np.number).columns.tolist()

plt.figure(figsize=(20,30))

for i, variable in enumerate(numerical_col):
    plt.subplot(5,4,i+1)
    plt.boxplot(data[variable],whis=1.5)
    plt.tight_layout()
    plt.title(variable)

plt.show()
```



- There are no outliers now

## Transform Age into bigger groups

```
In [82]: def age_grouping(age):
         return (age//10)*10
```

```
In [83]: data['AgeGroup'] = data['Age'].apply(age_grouping)

         data.AgeGroup.unique()
```


```
Out[83]: array([40., 30., 50., 20., 10., 60.])
```

```
In [84]: data = data.drop(['Age'], axis=1)
```

```
In [85]: print('AgeGroup\n' , data['AgeGroup'].value_counts(normalize=True) , '\n')
         bar_count_pct(data.AgeGroup)
```

```
AgeGroup
30.0    0.450900
40.0    0.231792
20.0    0.177987
50.0    0.122136
10.0    0.009411
60.0    0.007774
Name: AgeGroup, dtype: float64
```

```
Top:30.0
Freq:2204
```



- Now it is more clear as to what age group is most prevalent in the dataset.
- Finally, let convert ProdTaken to category so it is handled correctly

```
In [86]: data['ProdTaken'] = data.ProdTaken.astype('category')
```

```
In [87]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4888 entries, 0 to 4887
Data columns (total 19 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   ProdTaken                            4888 non-null   category
1   TypeofContact                        4888 non-null   category
2   CityTier                             4888 non-null   category
3   DurationOfPitch                      4888 non-null   float64
4   Occupation                           4888 non-null   category
5   Gender                               4888 non-null   object
6   NumberOfPersonVisiting               4888 non-null   float64
7   NumberOfFollowups                    4888 non-null   float64
8   ProductPitched                       4888 non-null   category
9   PreferredPropertyStar                4888 non-null   category
10  MaritalStatus                        4888 non-null   category
11  NumberOfTrips                        4888 non-null   float64
12  Passport                             4888 non-null   category
13  PitchSatisfactionScore                4888 non-null   category
14  OwnCar                               4888 non-null   category
15  NumberOfChildrenVisiting             4888 non-null   float64
16  Designation                          4888 non-null   category
17  MonthlyIncome                        4888 non-null   float64
18  AgeGroup                             4888 non-null   float64
```

```
dtypes: category(11), float64(7), object(1)
memory usage: 359.7+ KB
```

In [ ]:

# Model building - Bagging

## Split the Data

- Because there is a significant imbalance in the distribution of the target classes (18% success on ProdTaken vs 82%), we will use stratified sampling to ensure that relative class frequencies are approximately preserved in train and test sets.
- We will be using the stratify parameter in the train\_test\_split function.

```
In [88]: #X = creditData.drop("default" , axis=1)
#y = creditData.pop("default")

X = data.drop(['ProdTaken'],axis=1)
X = pd.get_dummies(X,drop_first=True)
#y = data['ProdTaken'].apply(lambda x : 1 if x=='Yes' else 0)
y = data.pop('ProdTaken')
```

```
In [89]: #X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.30, random_state=

# Splitting data into training and test set:
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1,
print(X_train.shape, X_test.shape)

(3421, 34) (1467, 34)
```

## Create Functions

```
In [90]: ## Function to calculate recall score
def get_recall_score(model,flag=True):
    """
    model : classifier to predict values of X

    """
    a = [] # defining an empty list to store train and test results
    pred_train = model.predict(X_train)
    pred_test = model.predict(X_test)
    train_recall = metrics.recall_score(y_train,pred_train)
    test_recall = metrics.recall_score(y_test,pred_test)
    a.append(train_recall) # adding train recall to list
    a.append(test_recall) # adding test recall to list
    if flag == True: # If the flag is set to True then only the following print statement
        print("Recall on training set : ",metrics.recall_score(y_train,pred_train))
        print("Recall on test set : ",metrics.recall_score(y_test,pred_test))

    return a # returning the list with train and test scores
```

```
In [91]: ## Function to calculate precision score
def get_precision_score(model,flag=True):
    """
    model : classifier to predict values of X
```

```

'''
b = [] # defining an empty list to store train and test results
pred_train = model.predict(X_train)
pred_test = model.predict(X_test)
train_precision = metrics.precision_score(y_train,pred_train)
test_precision = metrics.precision_score(y_test,pred_test)
b.append(train_precision) # adding train precision to list
b.append(test_precision) # adding test precision to list
if flag == True: # If the flag is set to True then only the following print stateme
    print("Precision on training set : ",metrics.precision_score(y_train,pred_train)
    print("Precision on test set : ",metrics.precision_score(y_test,pred_test))

return b # returning the list with train and test scores

```

```

In [92]: ## Function to calculate accuracy score
def get_accuracy_score(model,flag=True):
    '''
    model : classifier to predict values of X

    '''

    c = [] # defining an empty list to store train and test results
    train_acc = model.score(X_train,y_train)
    test_acc = model.score(X_test,y_test)
    c.append(train_acc) # adding train accuracy to list
    c.append(test_acc) # adding test accuracy to list
    if flag == True: # If the flag is set to True then only the following print stateme
        print("Accuracy on training set : ",model.score(X_train,y_train))
        print("Accuracy on test set : ",model.score(X_test,y_test))

    return c # returning the list with train and test scores

```

```

In [93]: def make_confusion_matrix(model,y_actual,labels=[1, 0]):
    '''
    model : classifier to predict values of X
    y_actual : ground truth

    '''

    y_predict = model.predict(X_test)
    cm=metrics.confusion_matrix( y_actual, y_predict, labels=[0, 1])
    df_cm = pd.DataFrame(cm, index = [i for i in ["Actual - No","Actual - Yes"]],
        columns = [i for i in ['Predicted - No','Predicted - Yes']])
    group_counts = ["{0:0.0f}".format(value) for value in
        cm.flatten()]
    group_percentages = ["{0:.2%}".format(value) for value in
        cm.flatten()/np.sum(cm)]
    labels = [f"{v1}\n{v2}" for v1, v2 in
        zip(group_counts,group_percentages)]
    labels = np.asarray(labels).reshape(2,2)
    plt.figure(figsize = (10,7))
    sns.heatmap(df_cm, annot=labels,fmt='')
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

```

## Build Decision Tree Model

- We will build our model using the DecisionTreeClassifier function. Using default 'gini' criteria to split.




- We will pass a dictionary {0:0.18,1:0.82} to the model to specify the weight of each class and the decision tree will give more weightage to class 1.

```
In [94]: dtree = DecisionTreeClassifier(criterion='gini',class_weight={0:0.18,1:0.82},random_sta
```

```
In [95]: dtree.fit(X_train, y_train)
```

```
Out[95]: DecisionTreeClassifier(class_weight={0: 0.18, 1: 0.82}, random_state=1)
```

```
In [96]: make_confusion_matrix(dtree,y_test)
```



Confusion Matrix -

- Consumer took the product and the model predicted it correctly that ProdTaken=1 : True Positive (observed=1,predicted=1)
- Consumer didn't take the product and the model predicted ProdTaken=1 : False Positive (observed=0,predicted=1)
- Consumer didn't take the product and the model predicted ProdTaken=0 : True Negative (observed=0,predicted=0)
- Consumer took the product and the model predicted that ProdTaken=0 : False Negative (observed=1,predicted=0)

```
In [97]: dtree_acc = get_accuracy_score(dtree)
dtree_recall = get_recall_score(dtree)
dtree_precision = get_precision_score(dtree)
```

Accuracy on training set : 1.0  
Accuracy on test set : 0.89093387866394  
Recall on training set : 1.0

Recall on test set : 0.7318840579710145  
Precision on training set : 1.0  
Precision on test set : 0.7013888888888888


- Decision Treemodel is overfitting on the training set and is performing poorly on the test set in terms of recall.

## Build Bagging Classifier Model

```
In [98]: bagging = BaggingClassifier(random_state=1)  
bagging.fit(X_train,y_train)
```

```
Out[98]: BaggingClassifier(random_state=1)
```

```
In [99]: make_confusion_matrix(bagging,y_test)
```



```
In [100... bagging_acc = get_accuracy_score(bagging)  
bagging_recall = get_recall_score(bagging)  
bagging_precision = get_precision_score(bagging)
```

Accuracy on training set : 0.9944460684010523  
Accuracy on test set : 0.9107021131561008  
Recall on training set : 0.9720496894409938  
Recall on test set : 0.644927536231884  
Precision on training set : 0.9984051036682615  
Precision on test set : 0.8436018957345972


- Bagging classifier is overfitting on the training set and is performing poorly on the test set in terms of recall.

## Bagging Classifier with weighted decision tree

```
In [101...] bagging_wt = BaggingClassifier(base_estimator=DecisionTreeClassifier(criterion='gini', c
bagging_wt.fit(X_train,y_train)
```

```
Out[101...] BaggingClassifier(base_estimator=DecisionTreeClassifier(class_weight={0: 0.18,
                                                                    1: 0.82},
                                                                    random_state=1),
                                                                    random_state=1)
```

```
In [102...] make_confusion_matrix(bagging_wt,y_test)
```



```
In [103...] wt_bagging_acc = get_accuracy_score(bagging_wt)
wt_bagging_recall = get_recall_score(bagging_wt)
wt_bagging_precision = get_precision_score(bagging_wt)
```

Accuracy on training set : 0.9956153171587255  
Accuracy on test set : 0.9038854805725971  
Recall on training set : 0.9782608695652174  
Recall on test set : 0.5652173913043478  
Precision on training set : 0.9984152139461173  
Precision on test set : 0.8813559322033898


- Bagging classifier with a weighted decision tree is giving very good accuracy and prediction but is not able to generalize well on test data in terms of recall.

## Build Random Forest Model

```
In [104...] rf = RandomForestClassifier(random_state=1)
rf.fit(X_train,y_train)
```

Out[104... RandomForestClassifier(random\_state=1)

```
In [105... make_confusion_matrix(rf,y_test)
```



```
In [106... rf_acc = get_accuracy_score(rf)
rf_recall = get_recall_score(rf)
rf_precision = get_precision_score(rf)
```

Accuracy on training set : 1.0  
Accuracy on test set : 0.9141104294478528  
Recall on training set : 1.0  
Recall on test set : 0.5797101449275363  
Precision on training set : 1.0  
Precision on test set : 0.9411764705882353


- Random Forest has performed well in terms of accuracy and precision, but it is not able to generalize well on the test data in terms of recall.

## Random forest with class weights

```
In [107... rf_wt = RandomForestClassifier(class_weight={0:0.18,1:0.82}, random_state=1)
rf_wt.fit(X_train,y_train)
```

Out[107... RandomForestClassifier(class\_weight={0: 0.18, 1: 0.82}, random\_state=1)

```
In [108... make_confusion_matrix(rf_wt,y_test)
```



```
In [109... wt_rf_acc = get_accuracy_score(rf_wt)
wt_rf_recall = get_recall_score(rf_wt)
wt_rf_precision = get_precision_score(rf_wt)
```

```
Accuracy on training set : 1.0
Accuracy on test set : 0.901840490797546
Recall on training set : 1.0
Recall on test set : 0.5253623188405797
Precision on training set : 1.0
Precision on test set : 0.9177215189873418
```

- There is not much improvement in metrics of weighted random forest as compared to the unweighted random forest.
- Random Forest with Weighted Tree has performed well in terms of accuracy and precision, but it is not able to generalize well on the test data in terms of recall.

## Tuning Models

### Using GridSearch for Hyperparameter tuning model

- Grid search is a tuning technique that attempts to compute the optimum values of hyperparameters.
- It is an exhaustive search that is performed on a the specific parameter values of a model.

### Tuning Decision Tree

```
In [110... # Choose the type of classifier.
dtree_estimator = DecisionTreeClassifier(class_weight={0:0.18,1:0.82},random_state=1)
```

```

# Grid of parameters to choose from - this will perform 2800 tests
parameters = {'max_depth': np.arange(2,30),
              'min_samples_leaf': [1, 2, 5, 7, 10],
              'max_leaf_nodes' : [2, 3, 5, 10,15],
              'min_impurity_decrease': [0.0001,0.001,0.01,0.1]
              }

# Type of scoring used to compare parameter combinations
scorer = metrics.make_scorer(metrics.recall_score)

# Run the grid search
grid_obj = GridSearchCV(dtree_estimator, parameters, scoring=scorer)
grid_obj = grid_obj.fit(X_train, y_train)


# Set the clf to the best combination of parameters
dtree_estimator = grid_obj.best_estimator_

# Fit the best algorithm to the data.
dtree_estimator.fit(X_train, y_train)

```

Out[110...] DecisionTreeClassifier(class\_weight={0: 0.18, 1: 0.82}, max\_depth=2, max\_leaf\_nodes=2, min\_impurity\_decrease=0.1, random\_state=1)

In [111...] make\_confusion\_matrix(dtree\_estimator,y\_test)



In [112...] tuned\_dtree\_acc = get\_accuracy\_score(dtree\_estimator)  
tuned\_dtree\_recall = get\_recall\_score(dtree\_estimator)  
tuned\_dtree\_precision = get\_precision\_score(dtree\_estimator)

Accuracy on training set : 0.1882490499853844  
Accuracy on test set : 0.18813905930470348  
Recall on training set : 1.0  
Recall on test set : 1.0

Precision on training set : 0.1882490499853844  
Precision on test set : 0.18813905930470348

- Overfitting in decision tree has reduced Accuracy and Precision, but the Recall has improved.  
This is an indication that overall the model is making many mistakes.

## Tuning Bagging Classifier

```
In [113... # grid search for bagging classifier
c11 = DecisionTreeClassifier(class_weight={0:0.18,1:0.82},random_state=1)
param_grid = {'base_estimator':[c11],
              'n_estimators':[5,7,15,51,101],
              'max_features': [0.7,0.8,0.9,1]
              }


grid = GridSearchCV(BaggingClassifier(random_state=1,bootstrap=True), param_grid=param_
grid.fit(X_train, y_train)
```

```
Out[113... GridSearchCV(cv=5, estimator=BaggingClassifier(random_state=1),
              param_grid={'base_estimator': [DecisionTreeClassifier(class_weight={0: 0.1
8,
                                                                    1: 0.8
2},
                                                                    random_state=1)],
              'max_features': [0.7, 0.8, 0.9, 1],
              'n_estimators': [5, 7, 15, 51, 101]},
              scoring='recall')
```

```
In [114... ## getting the best estimator
bagging_estimator = grid.best_estimator_
bagging_estimator.fit(X_train,y_train)
```

```
Out[114... BaggingClassifier(base_estimator=DecisionTreeClassifier(class_weight={0: 0.18,
                                                                    1: 0.82},
                                                                    random_state=1),
              max_features=1, n_estimators=51, random_state=1)
```

```
In [115... make_confusion_matrix(bagging_estimator,y_test)
```



```
In [116... tuned_bagging_acc= get_accuracy_score(bagging_estimator)
tuned_bagging_recall = get_recall_score(bagging_estimator)
tuned_bagging_precision = get_precision_score(bagging_estimator)
```

Accuracy on training set : 0.4811458637825197  
Accuracy on test set : 0.4226312201772324  
Recall on training set : 0.9611801242236024  
Recall on test set : 0.8913043478260869  
Precision on training set : 0.2612916842549599  
Precision on test set : 0.23142050799623706

- Recall has improved but the accuracy and precision of the model has dropped drastically which is an indication that overall the model is making many mistakes.

## Tuning Random Forest

```
In [117... # Choose the type of classifier.
rf_estimator = RandomForestClassifier(random_state=1)

# Grid of parameters to choose from
parameters = {
    "n_estimators": [110,251],
    "min_samples_leaf": np.arange(1, 6,1),
    "max_features": [0.7,0.9,'log2','auto'],
    "max_samples": [0.7,0.9,None],
}

# Run the grid search
grid_obj = GridSearchCV(rf_estimator, parameters, scoring='recall',cv=5)
grid_obj = grid_obj.fit(X_train, y_train)
```



```
# Set the clf to the best combination of parameters
rf_estimator = grid_obj.best_estimator_

# Fit the best algorithm to the data.
rf_estimator.fit(X_train, y_train)
```

Out[117... RandomForestClassifier(max\_features=0.7, n\_estimators=251, random\_state=1)

In [118... make\_confusion\_matrix(rf\_estimator,y\_test)



In [119... tuned\_rf\_acc = get\_accuracy\_score(rf\_estimator)  
tuned\_rf\_recall = get\_recall\_score(rf\_estimator)  
tuned\_rf\_precision = get\_precision\_score(rf\_estimator)

Accuracy on training set : 1.0  
Accuracy on test set : 0.929107021131561  
Recall on training set : 1.0  
Recall on test set : 0.7028985507246377  
Precision on training set : 1.0  
Precision on test set : 0.8981481481481481

- Recall has improved and the accuracy and precision of the model has improved drastically.

## Comparing all the models

In [122... # defining List of models  
models = [dtree,dtree\_estimator,bagging,bagging\_wt,bagging\_estimator,rf,rf\_wt,rf\_estima  
# defining empty lists to add train and test results  
acc\_train = []  
acc\_test = []  
recall\_train = []

```

recall_test = []
precision_train = []
precision_test = []

# Looping through all the models to get the accuracy, recall and precision scores
for model in models:
    # accuracy score
    j = get_accuracy_score(model, False)
    acc_train.append(j[0])
    acc_test.append(j[1])
    # recall score
    k = get_recall_score(model, False)
    recall_train.append(k[0])
    recall_test.append(k[1])
    # precision score
    l = get_precision_score(model, False)
    precision_train.append(l[0])
    precision_test.append(l[1])

```

```

In [123...] comparison_frame = pd.DataFrame({'Model': ['Decision Tree', 'Tuned Decision Tree', 'Bagging
                                                    'Weighted Bagging Classifier', 'Tuned Bagging
                                                    'Random Forest', 'Weighted Random Forest', 'Tun
                                                    'Train_Accuracy': acc_train,
                                                    'Test_Accuracy': acc_test,
                                                    'Train_Recall': recall_train,
                                                    'Test_Recall': recall_test,
                                                    'Train_Precision': precision_train,
                                                    'Test_Precision': precision_test})

comparison_frame

```

```

Out[123...]

```

	Model	Train_Accuracy	Test_Accuracy	Train_Recall	Test_Recall	Train_Precision	Test_Precision
0	Decision Tree	1.000000	0.890934	1.000000	0.731884	1.000000	0.701389
1	Tuned Decision Tree	0.188249	0.188139	1.000000	1.000000	0.188249	0.188139
2	Bagging Classifier	0.994446	0.910702	0.972050	0.644928	0.998405	0.843602
3	Weighted Bagging Classifier	0.995615	0.903885	0.978261	0.565217	0.998415	0.881356
4	Tuned Bagging Classifier	0.481146	0.422631	0.961180	0.891304	0.261292	0.231421
5	Random Forest	1.000000	0.914110	1.000000	0.579710	1.000000	0.941176
6	Weighted Random Forest	1.000000	0.901840	1.000000	0.525362	1.000000	0.917722
7	Tuned Random Forest	1.000000	0.929107	1.000000	0.702899	1.000000	0.898148

- Decision tree performed well on training and test set.
- Bagging classifier overfitted the data before and after tuning.
- Random Forest with default parameters performed better after tuning.

## Feature importance of Random Forest with Tuning

```
In [124... # importance of features in the tree building ( The importance of a feature is computed
#(normalized) total reduction of the criterion brought by that feature. It is also know


print (pd.DataFrame(rf_estimator.feature_importances_, columns = ["Imp"], index = X_tra
```

	Imp
MonthlyIncome	0.149993
DurationOfPitch	0.118809
Passport_1	0.067302
NumberOfTrips	0.065421
NumberOfFollowups	0.059158
AgeGroup	0.055359
Designation_Executive	0.050722
CityTier_3	0.038826
MaritalStatus_Single	0.032236
NumberOfPersonVisiting	0.028049
NumberOfChildrenVisiting	0.027526
Gender_Male	0.026446
PreferredPropertyStar_5.0	0.024879
PitchSatisfactionScore_3	0.022245
TypeofContact_Self Enquiry	0.022170
MaritalStatus_Unmarried	0.020512
PitchSatisfactionScore_4	0.019415
PreferredPropertyStar_4.0	0.018364
OwnCar_1	0.018127
Occupation_Large Business	0.017021
PitchSatisfactionScore_5	0.016978
Occupation_Small Business	0.015879
Occupation_Salaried	0.014933
MaritalStatus_Married	0.014003
CityTier_2	0.011154
ProductPitched_Deluxe	0.008305
Designation_Manager	0.007794
PitchSatisfactionScore_2	0.007411
ProductPitched_Standard	0.006327
Designation_Senior Manager	0.006286
ProductPitched_Super Deluxe	0.003037
ProductPitched_King	0.002246
Designation_VP	0.002162
TypeofContact_Unknown	0.000905

```
In [125... feature_names = X_train.columns
```

```
In [126... importances = rf_estimator.feature_importances_
indices = np.argsort(importances)

plt.figure(figsize=(12,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='violet', align='center')
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



- Monthlyincome is the most important feature for prediction followed by DurationOfPitch , Passport\_1 and NumberOfTrips .

In [ ]:

In [ ]:

## Model building - Boosting

In [135...]

```
## Function to calculate r2_score and RMSE on train and test data
def get_model_score(model, flag=True):
    """
    model : classifier to predict values of X

    """
    # defining an empty list to store train and test results
    score_list=[]

    pred_train = model.predict(X_train)
    pred_test = model.predict(X_test)

    train_r2=metrics.r2_score(y_train,pred_train)
```

```

test_r2=metrics.r2_score(y_test,pred_test)
train_rmse=np.sqrt(metrics.mean_squared_error(y_train,pred_train))
test_rmse=np.sqrt(metrics.mean_squared_error(y_test,pred_test))

#Adding all scores in the list
score_list.extend((train_r2,test_r2,train_rmse,test_rmse))

# If the flag is set to True then only the following print statements will be displayed
if flag==True:
    print("R-square on training set : ",metrics.r2_score(y_train,pred_train))
    print("R-square on test set : ",metrics.r2_score(y_test,pred_test))
    print("RMSE on training set : ",np.sqrt(metrics.mean_squared_error(y_train,pred_train)))
    print("RMSE on test set : ",np.sqrt(metrics.mean_squared_error(y_test,pred_test)))

# returning the list with train and test scores
return score_list

```

## AdaBoost Regressor Model

```

In [148... abc = AdaBoostClassifier(random_state=1)
abc.fit(X_train,y_train)

```

```


Out[148... AdaBoostClassifier(random_state=1)

```

```

In [149... make_confusion_matrix(abc,y_test)

```



```

In [147... abc_score=get_metrics_score(abc)

Accuracy on training set : 0.842443729903537
Accuracy on test set : 0.841854124062713
Recall on training set : 0.3059006211180124
Recall on test set : 0.30434782608695654

```

Precision on training set : 0.6816608996539792  
Precision on test set : 0.6774193548387096

- AdaBoost is generalizing well but it is giving very poor performance on recall.

## Hyperparameter Tuning

```
In [150... # Choose the type of classifier.
abc_tuned = AdaBoostClassifier(random_state=1)

# Grid of parameters to choose from
## add from article
parameters = {
    #Let's try different max_depth for base_estimator
    "base_estimator": [DecisionTreeClassifier(max_depth=1), DecisionTreeClassifier(max_de
    "n_estimators": np.arange(10, 110, 10),
    "learning_rate": np.arange(0.1, 2, 0.1)
}

# Type of scoring used to compare parameter combinations
acc_scorer = metrics.make_scorer(metrics.recall_score)


# Run the grid search
grid_obj = GridSearchCV(abc_tuned, parameters, scoring=acc_scorer, cv=5)
grid_obj = grid_obj.fit(X_train, y_train)

# Set the clf to the best combination of parameters
abc_tuned = grid_obj.best_estimator_

# Fit the best algorithm to the data.
abc_tuned.fit(X_train, y_train)
```

```
Out[150... AdaBoostClassifier(base_estimator=DecisionTreeClassifier(max_depth=3),
                           learning_rate=1.5000000000000002, n_estimators=100,
                           random_state=1)
```

```
In [151... make_confusion_matrix(abc_tuned, y_test)
```




```
In [152... abc_tuned_score=get_metrics_score(abc_tuned)
```

```
Accuracy on training set : 0.9874305758550131
Accuracy on test set : 0.8752556237218814
Recall on training set : 0.953416149068323
Recall on test set : 0.6159420289855072
Precision on training set : 0.9792663476874003
Precision on test set : 0.6882591093117408
```

- The model is overfitting the train data as train accuracy is much higher than the test accuracy.
- The model has low test recall. This implies that the model is not good at identifying defaulters.

```
In [153... importances = abc_tuned.feature_importances_
indices = np.argsort(importances)
feature_names = list(X.columns)

plt.figure(figsize=(12,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='violet', align='center')
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



- MonthlyIncome is the most important feature as per the tuned AdaBoost model, followed by DurationOfPitch , NumberOfTrips and AgeGroup .


## Gradient Boosting Classifier Model

```
In [154... gbc_init = GradientBoostingClassifier(init=AdaBoostClassifier(random_state=1),random_st
gbc_init.fit(X_train,y_train)
```

```
Out[154... GradientBoostingClassifier(init=AdaBoostClassifier(random_state=1),
random_state=1)
```

```
In [155... make_confusion_matrix(gbc_init,y_test)
```





```
In [156... #Using above defined function to get accuracy, recall and precision on train and test s
gbc_init_score=get_metrics_score(gbc_init)
```

Accuracy on training set : 0.8827828120432623  
 Accuracy on test set : 0.8663940013633266  
 Recall on training set : 0.44254658385093165  
 Recall on test set : 0.39855072463768115  
 Precision on training set : 0.8715596330275229  
 Precision on test set : 0.7857142857142857

- Gradient boosting is generalizing well and giving poor results on recall.

## Hyperparameter Tuning

```
In [157... # Choose the type of classifier.
gbc_tuned = GradientBoostingClassifier(init=AdaBoostClassifier(random_state=1),random_s

# Grid of parameters to choose from
## add from article
parameters = {
    "n_estimators": [100,150,200],
    "subsample": [0.8,0.9,1],
    "max_features": [0.7,0.8,0.9,1]
}

# Type of scoring used to compare parameter combinations
acc_scorer = metrics.make_scorer(metrics.recall_score)

# Run the grid search
grid_obj = GridSearchCV(gbc_tuned, parameters, scoring=acc_scorer,cv=5)
grid_obj = grid_obj.fit(X_train, y_train)


# Set the clf to the best combination of parameters
```

```
gbc_tuned = grid_obj.best_estimator_
```

```
# Fit the best algorithm to the data.  
gbc_tuned.fit(X_train, y_train)
```

```
Out[157...] GradientBoostingClassifier(init=AdaBoostClassifier(random_state=1),  
                                     max_features=0.9, n_estimators=200, random_state=1,  
                                     subsample=0.9)
```

```
In [159...] make_confusion_matrix(gbc_tuned,y_test)
```




```
In [160...] #Using above defined function to get accuracy, recall and precision on train and test s  
gbc_tuned_score=get_metrics_score(gbc_tuned)
```

```
Accuracy on training set : 0.9087985969014908  
Accuracy on test set : 0.8759372869802318  
Recall on training set : 0.562111801242236  
Recall on test set : 0.4855072463768116  
Precision on training set : 0.923469387755102  
Precision on test set : 0.7701149425287356
```

- The model performace has not increased by much.
- The model has started to overfit the train data in terms of recall.
- The model is generalizing well but it is giving very poor performance on recall.

```
In [161...] importances = gbc_tuned.feature_importances_  
indices = np.argsort(importances)  
feature_names = list(X.columns)  
  
plt.figure(figsize=(12,12))  
plt.title('Feature Importances')  
plt.barh(range(len(indices)), importances[indices], color='violet', align='center')
```

```
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



- `MonthlyIncome` is the most important feature as per the tuned AdaBoost model, followed by `Passport_1`, `Designation_Executive` and `DurationOfPitch`.

## XGBoost Classifier Model

```
In [165... # Choose the type of classifier.
xgb_tuned = XGBClassifier(random_state=1)

# Grid of parameters to choose from
## add from
parameters = {
    "n_estimators": np.arange(10,100,30),
    "scale_pos_weight": [0,1,2],
    "subsample": [0.5,0.7,1],
    "learning_rate": [0.01,0.1,0.2],
    "gamma": [0,1,3],
    "colsample_bytree": [0.5,0.7,1],
    "colsample_bylevel": [0.5,0.7,1]
}
```

```

# Type of scoring used to compare parameter combinations
acc_scorer = metrics.make_scorer(metrics.recall_score)

# Run the grid search
grid_obj = GridSearchCV(xgb_tuned, parameters,scoring=acc_scorer,cv=5)
grid_obj = grid_obj.fit(X_train, y_train)

# Set the clf to the best combination of parameters
xgb_tuned = grid_obj.best_estimator_

# Fit the best algorithm to the data.
xgb_tuned.fit(X_train, y_train)

```

```


Out[165... XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1,
                        colsample_bynode=1, colsample_bytree=1, gamma=1, gpu_id=-1,
                        importance_type='gain', interaction_constraints='',
                        learning_rate=0.2, max_delta_step=0, max_depth=6,
                        min_child_weight=1, missing=nan, monotone_constraints='()',
                        n_estimators=70, n_jobs=4, num_parallel_tree=1, random_state=1,
                        reg_alpha=0, reg_lambda=1, scale_pos_weight=2, subsample=1,
                        tree_method='exact', validate_parameters=1, verbosity=None)

```

```

In [166... make_confusion_matrix(xgb_tuned,y_test)

```



```

In [167... #Using above defined function to get accuracy, recall and precision on train and test
xgb_tuned_score=get_metrics_score(xgb_tuned)

```

```

Accuracy on training set : 0.9929845074539608
Accuracy on test set : 0.9222903885480572
Recall on training set : 0.9782608695652174
Recall on test set : 0.7391304347826086
Precision on training set : 0.984375
Precision on test set : 0.8292682926829268


```

- This model performs really well on all metrics. Even though Recall on test data seems a bit low, it is not that bad and all other metrics look good.

In [168...

```
importances = xgb_tuned.feature_importances_
indices = np.argsort(importances)
feature_names = list(X.columns)

plt.figure(figsize=(12,12))
plt.title('Feature Importances')
plt.barh(range(len(indices)), importances[indices], color='violet', align='center')
plt.yticks(range(len(indices)), [feature_names[i] for i in indices])
plt.xlabel('Relative Importance')
plt.show()
```



- Designation\_Executive is the most important feature for prediction followed by Passport\_1, MaritalStatus\_Single and CityTier\_3.

## Comparing all models

In [171...

```
# defining list of models
models = [abc, abc_tuned, gbc_init, gbc_tuned, xgb_tuned]

# defining empty lists to add train and test results
acc_train = []
```

```

acc_test = []
recall_train = []
recall_test = []
precision_train = []
precision_test = []

# Looping through all the models to get the accuracy, precall and precision scores
for model in models:
    j = get_metrics_score(model, False)
    acc_train.append(np.round(j[0], 2))
    acc_test.append(np.round(j[1], 2))
    recall_train.append(np.round(j[2], 2))
    recall_test.append(np.round(j[3], 2))
    precision_train.append(np.round(j[4], 2))
    precision_test.append(np.round(j[5], 2))

```

```

In [173...] comparison_frame = pd.DataFrame({'Model': ['AdaBoost with default paramters', 'AdaBoost
                                                    'Gradient Boosting with init=AdaBoost', 'Grad
                                                    'XGBoost Tuned'],
                                          'Train_Accuracy': acc_train, 'Test_Accuracy':
                                          'Train_Recall': recall_train, 'Test_Recall': re
                                          'Train_Precision': precision_train, 'Test_Prec

comparison_frame

```

Out[173...]

	Model	Train_Accuracy	Test_Accuracy	Train_Recall	Test_Recall	Train_Precision	Test_Precision
0	AdaBoost with default paramters	0.84	0.84	0.31	0.30	0.68	0.68
1	AdaBoost Tuned	0.99	0.88	0.95	0.62	0.98	0.69
2	Gradient Boosting with init=AdaBoost	0.88	0.87	0.44	0.40	0.87	0.79
3	Gradient Boosting Tuned	0.91	0.88	0.56	0.49	0.92	0.77
4	XGBoost Tuned	0.99	0.92	0.98	0.74	0.98	0.83

- Tuned XGBoost model is the best model here. It has really high performance metrics, and consistent recall values.

## Business Recommendations

- Company can focus on targeting customer with these strong important features:
  - Designation\_Executive
  - Passport\_1
  - MaritalStatus\_Single
  - CityTier\_3

In [ ]:

In [ ]:

In [ ]:

# End-of-File